



University of Pittsburgh

# CS 1541 Introduction

Technology Advances

Wonsun Ahn

Department of Computer Science

School of Computing and Information



# Technology Advances

---



University of Pittsburgh



# Advances in Technology


- Technology has been advancing at lightning speed
- Architecture and IT as a whole were beneficiaries
- Technology advance is summarized by *Moore's Law*
  - You probably heard of it at some point. Something about ...
  - "X doubles every 18-24 months at constant cost"
- Is X:
  - CPU performance?
  - CPU clock frequency?
  - Transistors per CPU chip?



# Moore's Law

## Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.




Data source: Wikipedia ([https://en.wikipedia.org/wiki/Transistor\\_count](https://en.wikipedia.org/wiki/Transistor_count))

The data visualization is available at [OurWorldinData.org](https://www.ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

# Miniaturization of Transistors



Data source: Radamson, H.H.; He, X.; Zhang, Q.; Liu, J.; Cui, H.; Xiang, J.; Kong, Z.; Xiong, W.; Li, J.; Gao, J.; Yang, H.; Gu, S.; Zhao, X.; Du, Y.; Yu, J.; Wang, G. Miniaturization of CMOS. *Micromachines* **2019**, *10*, 293.

- Moore's Law has been driven by transistor miniaturization
  - CPU chip area hasn't changed much





# Future of Moore's Law

- The semiconductor industry has produced roadmaps
  - Semiconductor Industry Association (SIA):  
1977~1997
  - International Technology Roadmap for Semiconductors (ITRS):  
1998~2016
  - International Roadmap for Devices and Systems (IRDS):  
2017~Present

## ■ IRDS Lithography Projection (2020)


Year of Production	2018	2020	2022	2025	2028	2031	2034
Technology Node (nm)	7	5	3	2.1	1.5	1.0	0.7

- Moore's Law will continue into foreseeable future
- IRDS does not project significant increase in CPU chip size
- Increases in transistors will come from *transistor density*



# IRDS isn't Perfect

- ITRS (predecessor of IRDS) has made corrections before



- After all, you are trying to predict the future
- But architects rely on the roadmap to design future processors

# Moore's Law and Performance




- Did Moore's Law result in higher performance CPUs?
- When you decide on a CPU for your desktop, what number(s) do you look at to see how fast it is?
- Best way: try running your favorite apps on it.
- Everyone has different favorite apps so instead publish
  - Results from running a suite of benchmark apps (e.g. SPEC CPU)
  - Several important components of performance



# Components of Execution Time



- Processor activity happens on clock “ticks” or cycles



- On each tick, bits flow through logic gates and are latched

- Execution time =  $\frac{\text{seconds}}{\text{program}}$

$$\begin{aligned}\frac{\text{seconds}}{\text{program}} &= \frac{\text{cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{cycle}} \\ &= \frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}\end{aligned}$$

# How to improve Execution Time



$$\frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

## ■ Reduce $\frac{\text{seconds}}{\text{cycle}}$ :

- Clock frequency =  $\frac{\text{cycles}}{\text{second}}$  = reverse of  $\frac{\text{seconds}}{\text{cycle}}$
- Higher clock frequency (GHz) leads to shorter exec time

## ■ Reduce $\frac{\text{cycles}}{\text{instruction}}$ :

- Also known as CPI (Cycles Per Instruction)
- IPC (Instructions Per Cycle) =  $\frac{\text{instructions}}{\text{cycles}}$  = reverse of  $\frac{\text{cycles}}{\text{instructions}}$
- Higher IPC leads to shorter execution time

## ■ Reduce $\frac{\text{instructions}}{\text{program}}$ :

- Less instructions leads to shorter execution time
- ISAs that do a lot of work with one instruction shortens time

# Moore's Law impacts two layers



- Did Moore's Law result in higher performance CPUs?
- Law impacts both architecture and physical layers

Instruction Set Architecture

**Processor Organization**

Computer  
Architecture

**Transistor Implementation**

Physical Layer

- Processor Organization: many more transistors to use in design
- Transistor Implementation: smaller, more efficient transistors



# Moore's Law Impact on Architecture

- So where did architects use all those transistors?
- Well, we will learn this throughout the semester 😊
  - Pipelining ← Improves frequency
  - Parallel execution
  - Branch prediction
  - Speculative execution
  - Memory caching

} Improves IPC
- Let's go on to impact on the physical layer for now




# Moore's Law Impact on Physical Layer

- Frequency = transistor speed ×  
number of transistors between clock ticks
  - Transistor speed is decided by the physical layer
  - Number of transistors between ticks is CPU design
  
- Given a CPU design, faster transistors → faster CPUs
  - Intel's tick-tock model staggers transistor / design improvements
  - Tick: new generation due to new technology node
  - Tock: new generation due to new CPU design
  
- So did Moore's Law result in faster transistors?
  - In other words, are smaller transistors faster?



# Speed of Transistors

- Transistors are like faucets:




- To make a transistor go fast, do one of the following:
  - Increase water pressure (supply voltage)
  - Change faucet (transistor) design
- Transistor design can be changed in two ways:
  - Increase pipe thickness (reduce channel resistance  $R$ )
  - Reduce bucket size (reduce capacitance  $C$ )
  - $T_{switch} \propto RC$  (transistor switch delay is proportional to  $RC$ )



# Smaller Transistors are Faster!




- Transistors are like faucets!



- For transistor dimension  $d$  (node size)
  - Capacitance  $C \propto d$
  - Channel resistance  $R \propto \text{Channel length } L \propto d$
  - $T_{\text{switch}} \propto RC \propto d^2$
- Given the same *supply voltage*, smaller is faster!
- Did Moore's Law enjoy faster and faster frequencies?



# Yes, for a while ...




Source: Computer Architecture, A Quantitative Approach (6<sup>th</sup> ed.) by John Hennessy and David Patterson, 2017

- Due in large part to improvement in transistor speed
  - CPU design (pipelining) contributed but we'll discuss later



# But not so much lately




Source: Computer Architecture, A Quantitative Approach (6<sup>th</sup> ed.) by John Hennessy and David Patterson, 2017

## ■ Suddenly around 2003, frequency scaling stops. Why?

- Improvements in transistor speed stopped.
- CPU design (pipelining) has met its limits.



# Dent in CPU Performance



Source: <https://preshing.com/20120208/a-look-back-at-single-threaded-cpu-performance/>

- This caused a big dent in CPU performance circa 2003
- Improvements henceforth mostly came from IPC



# Why did frequency scaling stop? TDP.

## ■ *TDP (Thermal Design Power):*

- Maximum power (heat) that CPU is designed to generate
- Capped by the amount of heat cooling system can handle
- Cooling system hasn't improved much over generations

## ■ CPU Power = $A * N * CFV^2$ **must be < TDP**

- A = Activity factor (% of transistors with activity)
- N = Number of transistors
- C = Capacitance
- F = Frequency
- V = Supply Voltage



## ■ What happens to each factor with Moore's Law?

# CPU Power (with Fixed Voltage)



- Change in CPU Power  $\propto A * N * CFV^2$ , with fixed V:
  - A = Activity factor
  - N = Number of transistors  $\propto 1/d^2$  ↑ ↑
  - C = Capacitance (size of bucket)  $\propto d$  ↓
  - F  $\propto 1/T_{\text{switch}} \propto 1/d^2$  ↑ ↑ (if supply voltage is kept constant)
  - V = Supply Voltage (water pressure)→ CPU Power  $\propto 1/d^3$  ↑ ↑ ↑ !
  
- Recipe used until 1990's until power started to matter
  - That meant F could not keep increasing while meeting TDP
  
- Q) So how did CPU frequency keep increasing up to 2003?





# CPU Power (with Dennard Scaling)

- Goal: to scale frequency while keeping power constant
- By reducing  $V$  proportional to  $d$ ,  
then change in CPU Power  $\propto A * N * CFV^2$  is:
  - $A$  = Activity factor
  - $N$  = Number of transistors  $\propto 1/d^2 \uparrow \uparrow$
  - $C$  = Capacitance  $\propto d \downarrow$
  - $F \propto 1/d \uparrow$  (Can scale by  $1/d$  while keeping power constant!)
  - $V$  = Supply Voltage ( $\propto d$ )  $\downarrow \rightarrow V^2 \downarrow \downarrow$
- **Dennard Scaling:** Above recipe for scaling up frequency, while reducing supply voltage to keep power constant




# Dennard Scaling and $V_{th}$

- So, it's that easy? Just reduce  $V$  until you meet TDP?
- No, it's not that simple ☹.
- Reducing  $V_{dd}$  (supply voltage) affects CPU operation
  - As  $V_{dd}$  is reduced, transistor becomes slower (lower pressure)
  - Eventually, CPU stops working altogether
- Transistors need redesigning to work at lower voltage
  - $V_{th}$  is the threshold voltage when switching happens
  - $V_{th}$  needs to be reduced along with  $V_{dd}$  to maintain speed

Speed is determined by  $V_{dd}$  if  $V_{th}$  is fixed



## ■ RC Charging Curve of $V_G$ (Gate Voltage)




- $V_{dd}$  is high  $\rightarrow V_G$  reaches  $V_{th}$  slowly at  $T_{on}$  (high pressure)
- $V_{dd}'$  is high  $\rightarrow V_G$  reaches  $V_{th}$  slowly at  $T_{on}'$  (low pressure)

# Speed is maintained **with lower $V_{th}$**




## ■ RC Charging Curve of $V_G$



- Speed ( $T_{on}$ ) is maintained while reducing  $V_{dd}$  to  $V_{dd}'$ ,  
**but only if  $V_{th}$  is also reduced to  $V_{th}'$**




# Dennard Scaling paid dividends for a while



- Aided by repeated reductions of  $V_{dd}$  and  $V_{th}$



# Then Dennard Scaling ended at 2003



- Main reason was that  $V_{th}$  could no longer be reduced






# Limits to Dropping $V_{th}$

## ■ Subthreshold leakage

- Transistor leaks current even when gate is off ( $V_G = 0$ )




- This leakage current translates to leakage power
- Leakage worsens when  $V_{th}$  is dropped



# Leakage Power across Generations

- Leakage power has increased across technology nodes



Source: L. Yan, Jiong Luo and N. K. Jha, "Joint dynamic voltage scaling and adaptive body biasing for heterogeneous distributed real-time embedded systems," in *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 24, no. 7, pp. 1030-1041, July 2005



# CPU Power must now include leakage

- Previous power calculation was incomplete
  - CPU power is the sum of both dynamic and leakage power
- $\text{Power}_{\text{CPU}} \propto \text{Power}_{\text{dynamic}} + \text{Power}_{\text{leakage}}$ 
  - $\text{Power}_{\text{dynamic}} \propto A * N * C V_{\text{dd}}^2$
  - $\text{Power}_{\text{leakage}} \propto N * V_{\text{dd}} * e^{-V_{\text{th}}}$
  - Leakage worsens *exponentially* when  $V_{\text{th}}$  is dropped
  - Catch-22: when dropping  $V_{\text{th}}$ ,  $\text{Power}_{\text{dynamic}}$  ↓ but  $\text{Power}_{\text{leakage}}$  ↑↑
- $V_{\text{th}}$  can't be reduced further, so  $V_{\text{dd}}$  can't be reduced
- Dennard Scaling relies on reducing  $V_{\text{dd}}$ , so it's the end

# CPU Power (End of Dennard Scaling)



- What happens to frequency without Dennard Scaling?
- $\text{Power}_{\text{dynamic}} (\propto A * N * CFV^2) + \text{Power}_{\text{leakage}} (\propto N * V * e^{-V_{\text{th}}})$ 
  - A = Activity factor
  - N = Number of transistors ( $\propto 1/d^2$ ) ↑ ↑
  - C = Capacitance ( $\propto d$ ) ↓
  - V = Supply Voltage  $\Leftrightarrow$  (Due to fixed  $V_{\text{th}}$ )
  - F = Frequency ???
- To offset N, you actually have to *decrease* F
- Otherwise, if you want to maintain F, must decrease N
  - That is, you cannot power on all the transistors at any given point
  - Dark silicon: situation where chip is only partially powered



# Free Ride is Over

- “Free” speed improvements from transistors is over
- Now it’s up to architects to improve performance
  - Moore’s Law is still alive (although slowing down)
  - Architects are flooded with extra transistors each generation
  - But it’s hard to even keep them powered without reducing F!
- Now is a good time to discuss technology constraints
  - Since we already mentioned a big one: TDP