

Multiprocessors and Caching

CS 1541

Wonsun Ahn

Two ways to use multiple processors

- **Distributed (Memory) System**

- Processors **do not share memory** (and by extension **data**)
- Processors exchange data through network messages
- Programming standards:
 - Message Passing Interface (MPI) – C/C++ API for exchanging messages
 - Ajax (Asynchronous JavaScript and XML) – API for web apps
- Data exchange protocols: TCP/IP, UDP/IP, JSON, XML...

- **Shared Memory System** (a.k.a. Multiprocessor System)

- Processors **share memory** (and by extension **data**)
- Programming standards:
 - Pthreads (POSIX threads), Java threads – APIs for threading
 - OpenMP – Compiler #pragma directives for parallelization
- **Cache coherence protocol**: protocol for exchanging data among caches
→ Just like Ethernet, caches are part of a larger network of caches

Shared Data Review

- What bad thing can happen when you have shared data?
- Dataraces!
 - You should have learned it in CS 449.
 - But if you didn't, don't worry I'll go over it.

Review: Datarace Example

```
int shared = 0;
void *add(void *unused) {
    for(int i=0; i < 1000000; i++) { shared++; }
    return NULL;
}
int main() {
    pthread_t t;
    // Child thread starts running add
    pthread_create(&t, NULL, add, NULL);
    // Main thread starts running add
    add(NULL);
    // Wait until child thread completes
    pthread_join(t, NULL);
    printf("shared=%d\n", shared);
    return 0;
}
```

bash-4.2\$./datarace

shared=1085894

bash-4.2\$./datarace

shared=1101173

bash-4.2\$./datarace


shared=1065494

Q) What do you expect from running this? Maybe `shared=2000000` ?

A) Nondeterministic result! Due to **datarace** on `shared`.


Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`




Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`




Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`




Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`




Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`




Review: Datarace Example

- When two threads do `shared++`; initially `shared = 1`



Review: Datarace Example


- When two threads do `shared++`; initially `shared = 1`



Review: Datarace Example

- Why did this occur in the first place?
- Because data was **replicated** to CPU registers and each worked on its own copy!

- When two threads do `shared++`; initially `shared = 1`



Review: Datarace Example

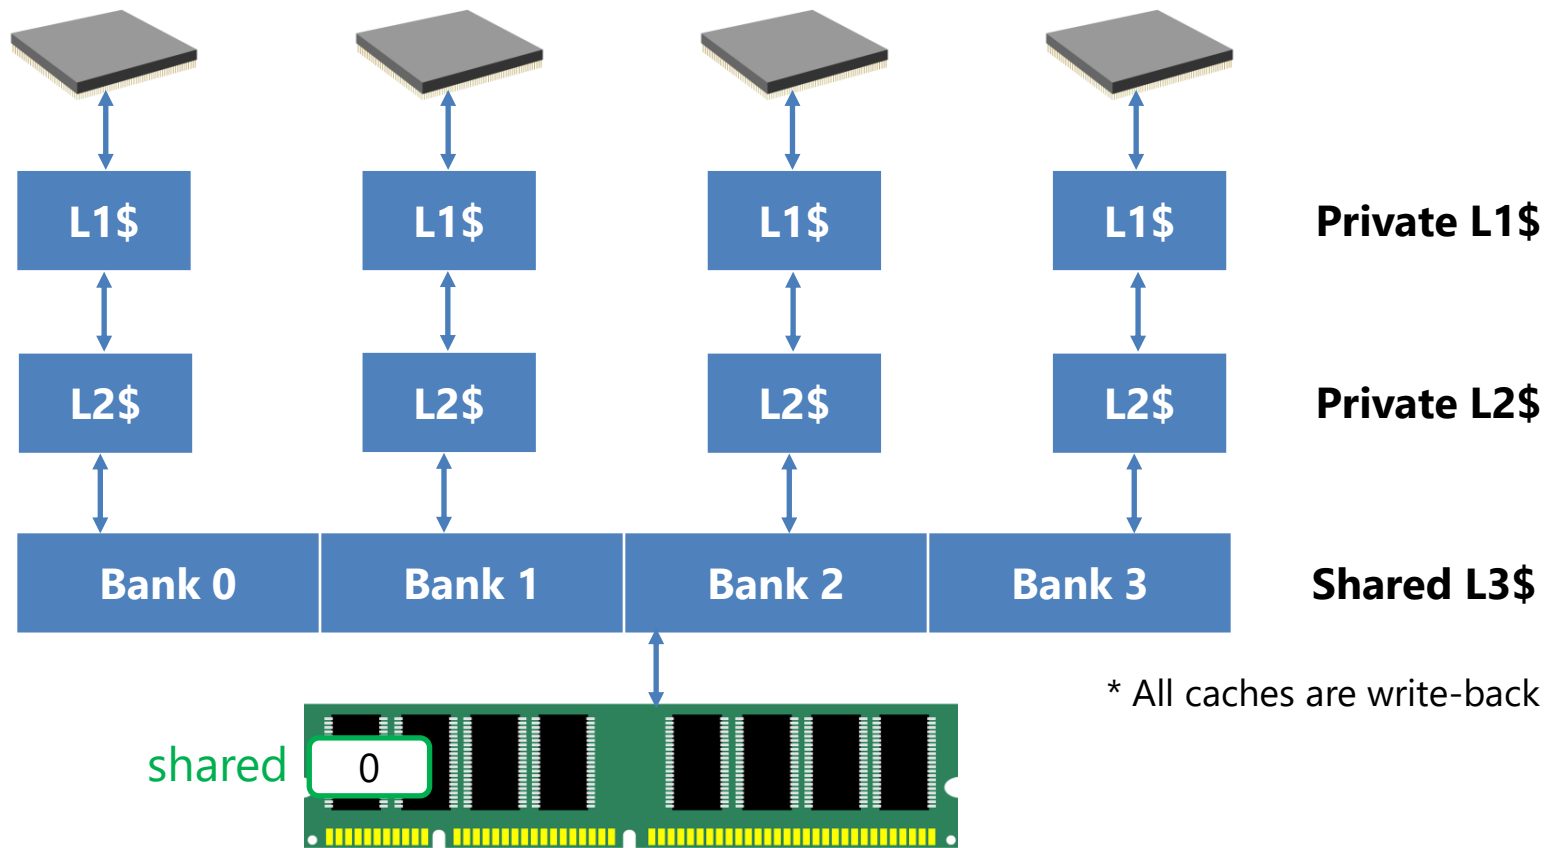
```
pthread_mutex_t lock;  
int shared = 0;  
void *add(void *unused) {  
    for(int i=0; i < 1000000; i++) {  
        pthread_mutex_lock(&lock);  
        shared++;  
        pthread_mutex_unlock(&lock);  
    }  
    return NULL;  
}  
int main() {  
    ...  
}
```

```
bash-4.2$ ./datarace  
shared=2000000  
bash-4.2$ ./datarace  
shared=2000000  
bash-4.2$ ./datarace  
shared=2000000
```

- Data race is fixed! Now shared is always 2000000.
- Problem solved? No! CPU registers is not the only place **replication** happens!


Caching also does replication!

- What happens if caches sit in between processors and memory?




Caching also does replication!

- Let's say CPU 0 first fetches `shared` for incrementing



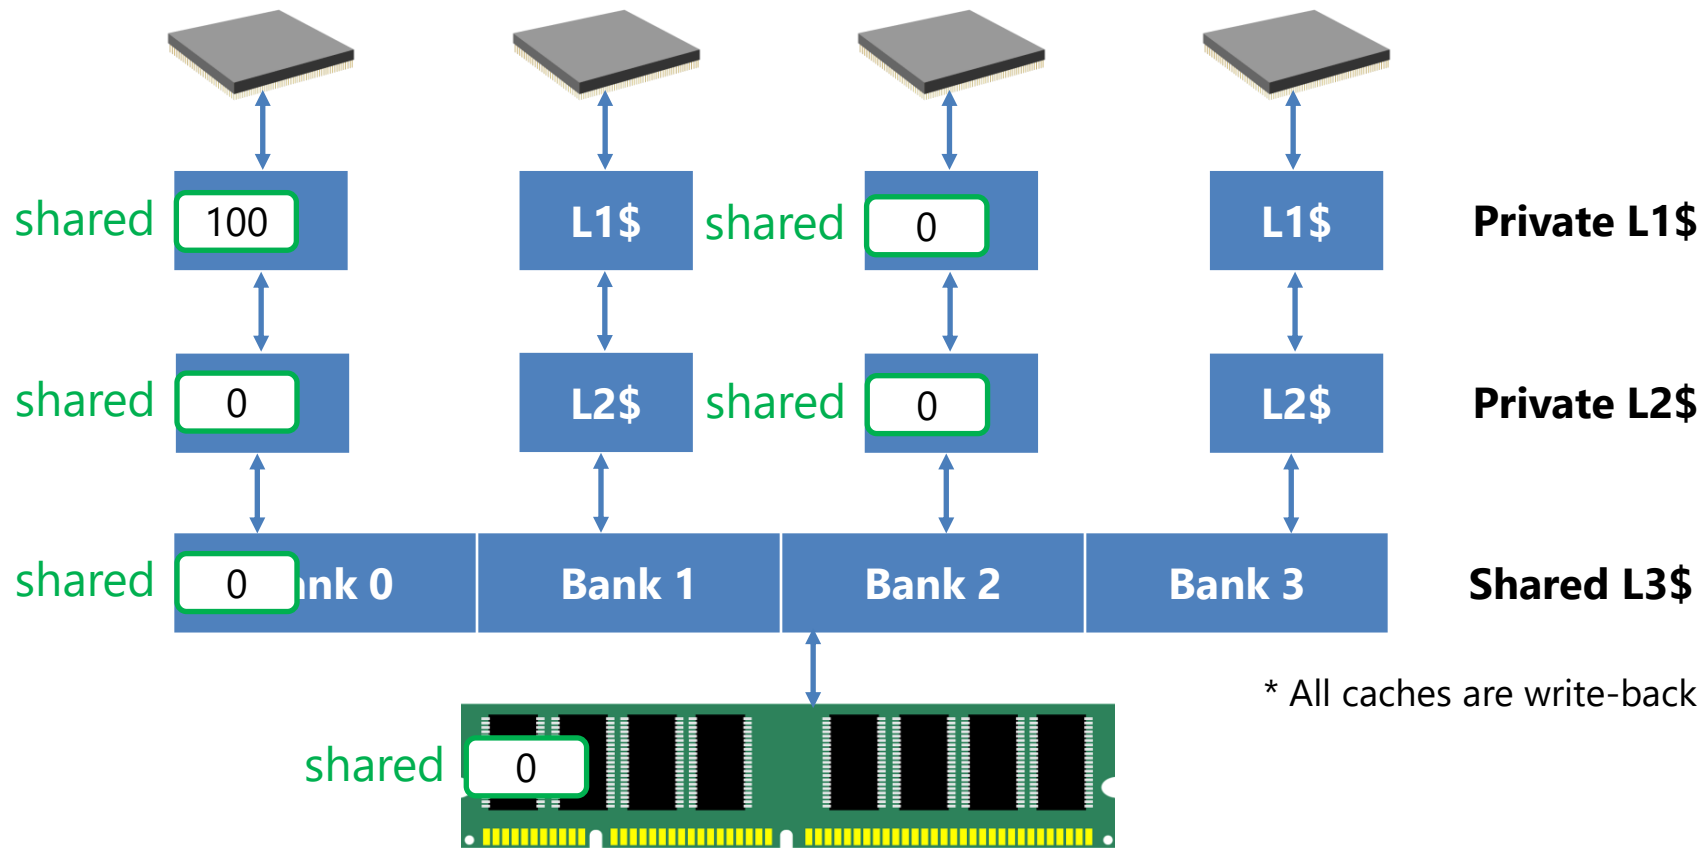
Caching also does replication!

- Then CPU 0 increments `shared` 100 times to 100



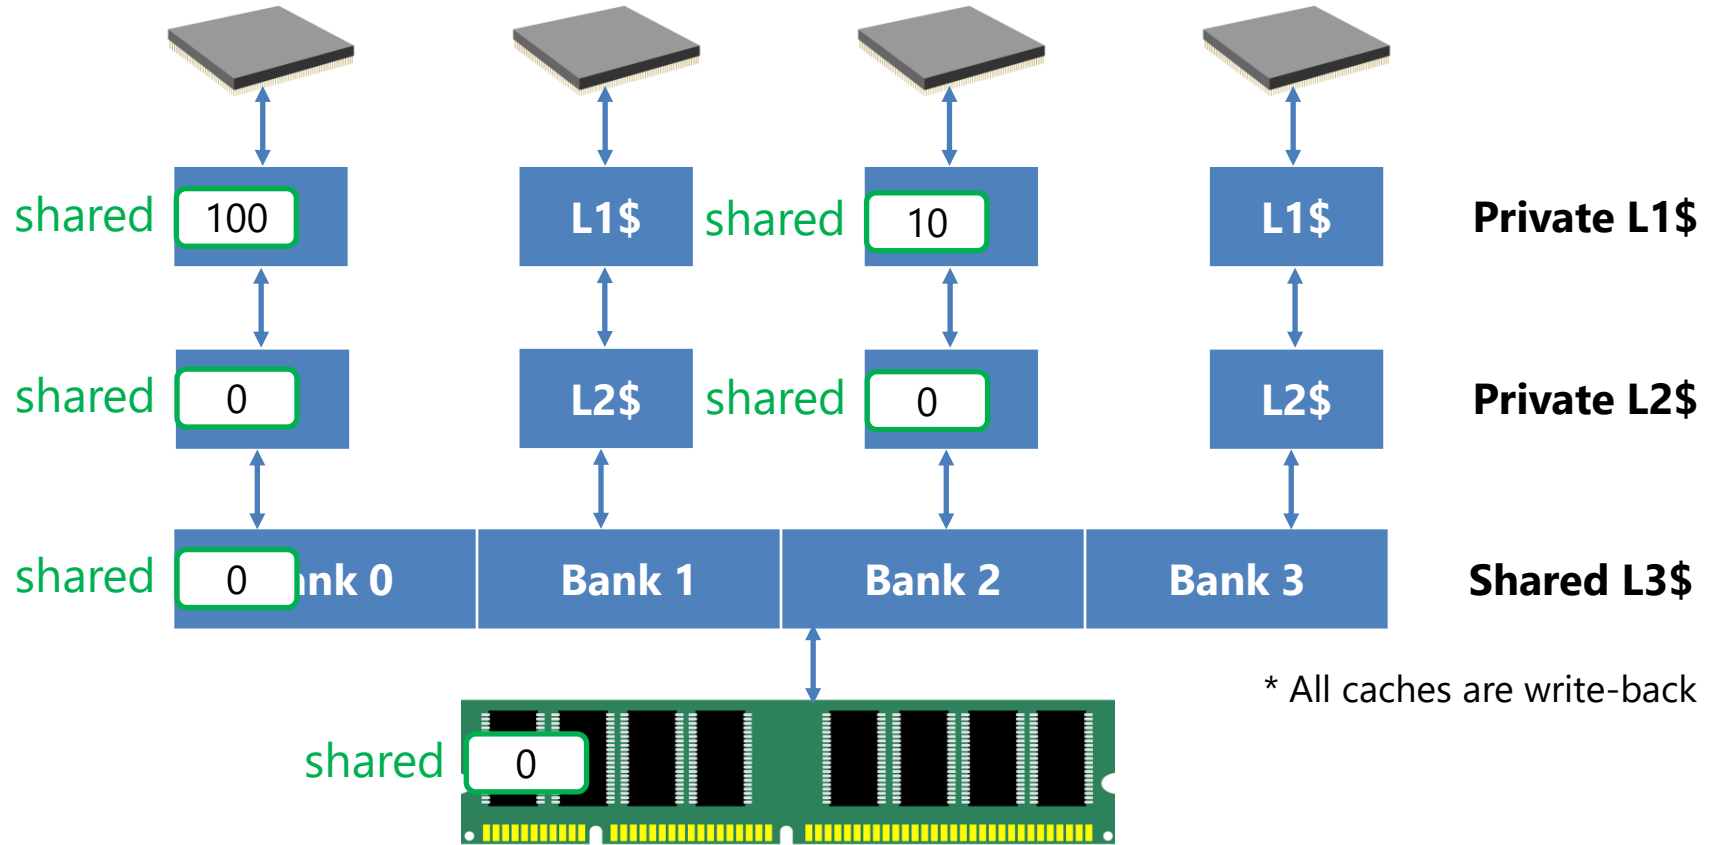
Caching also does replication!

- Then CPU 2 gets hold of the mutex and fetches `shared` from L3



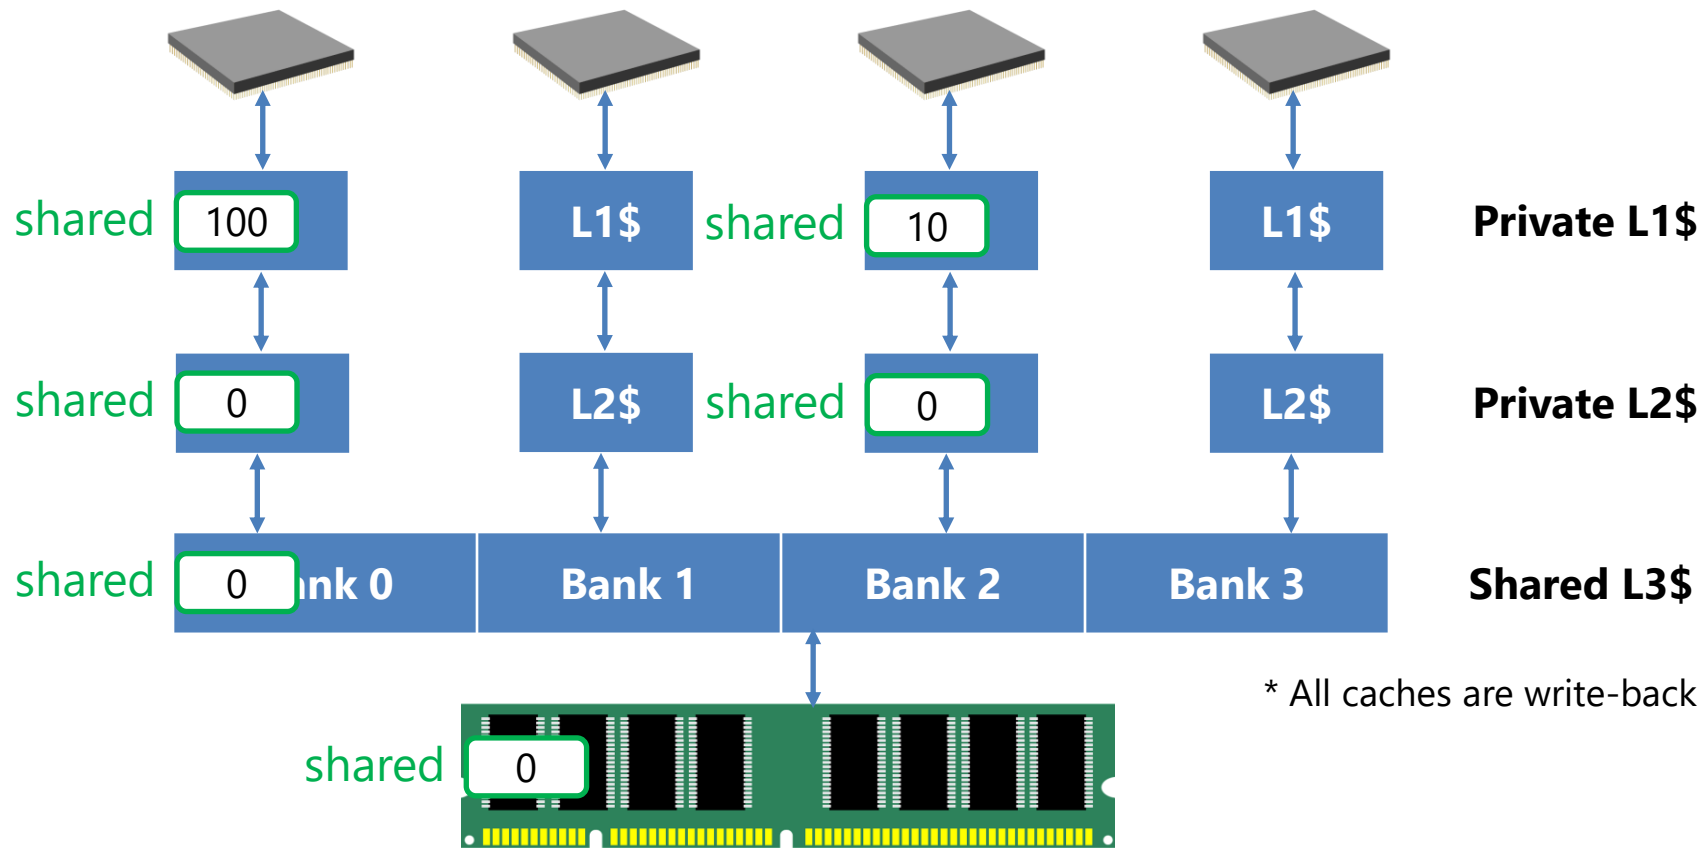
Caching also does replication!

- Then CPU 2 increments `shared` 10 times to 10



Caching also does replication!

- Clearly this is wrong. L1 caches of CPU 0 and CPU 2 are **incoherent**.



Cache Incoherence: Problem with Private Caches

- This problem does not occur with a **shared cache**.
 - All processors share and work on a **single copy** of data.



- The problem exists only with private caches.
- The problem exists for **private** caches.
 - Private copy is at times inconsistent with lower memory.
 - **Incoherence** occurs when private copies differ from each other.
 - Means processors return different values for same location!

Cache Coherence

Cache Coherence


- **Cache coherence** (loosely defined):
 - All processors of system should see the **same view of memory**
 - Copies of values cached by processors should adhere to this rule
- Each ISA has a different definition of what that “view” means
 - **Memory consistency model**: definition of what that “view” is
- All models agree on one thing:
 - That a change in value should reflect on all copies (eventually)

How Memory Consistent Model affects correctness

- Initially, `data == 0`, `flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```



Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```

Q) What do you expect the value of `data` will be when it gets printed?

A) Most people will say 42 because that is the logical ordering.

But is it? Not always. There are situations where `data` is still 0!

Scenario 1: Stores arrive out-of-order



- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```





Let's assume initially both `data` and `flag` are cached in each CPU's L1 caches.

Scenario 1: Stores arrive out-of-order

- Initially, `data == 0, flag == false`.



Producer Thread

```
data == 42;  
flag = true;
```



Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```





CPU 0 updates both `data` and `flag` to 42 and true.

Scenario 1: Stores arrive out-of-order

- Initially, `data == 0, flag == false`.



Producer Thread

```
data == 42;  
flag = true;
```



Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



Now the cached values in CPU 1 are stale and need to be **invalidated**.

Invalidation: act of marking a cache block with stale data invalid.

Scenario 1: Stores arrive out-of-order

- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



The invalidate messages travel through a network and may arrive out-of-order. Let's say invalidate for flag arrives first to CPU 1 and marks flag invalid.

Scenario 1: Stores arrive out-of-order

- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



CPU 1 fetches updated flag from CPU 0 when comparing `flag == false`.
Invalidate for data is still traveling through the network.

Scenario 1: Stores arrive out-of-order

- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```





Since `flag` is `true`, CPU 1 breaks out of while loop and prints `data`.
`data=0` gets printed!

Scenario 2: Loads perform out-of-order

- Initially, `data == 0, flag == false`.



Producer Thread

```
data == 42;  
flag = true;
```



Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



Let's assume now `flag` is not cached in CPU 1.



CPU 1 suffers a cache miss on `flag` when it compares `flag == false`.

Scenario 2: Loads perform out-of-order

- Initially, `data == 0, flag == false`.



Producer Thread

```
data == 42;  
flag = true;
```



Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



Instruction Queue

lw r1, flag **(miss)**

beq r1, \$zero, _loop

lw r2, data **(hit)**

call println on r2

Instead of stalling, CPU 1 predicts the branch not taken and issues `lw r2, data`.
Now, `r2 == 0`. (Unless pipeline flushes due to branch misprediction.)

Scenario 2: Loads perform out-of-order



- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



Fetch for flag

Instruction Queue

lw r1, flag (**miss**)

beq r1, \$zero, _loop

lw r2, data (**hit**)

call println on r2

Now let's say CPU 0 updates data and flag before the fetch for flag arrives.

Now, `lw r1, flag` completes, allowing `beq r1, $zero, _loop` to issue (with `r1 == true`)

Scenario 2: Loads perform out-of-order


- Initially, `data == 0, flag == false`.

Producer Thread


```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```



Fetch for flag



Instruction Queue

`lw r1, flag` (**miss**)

`beq r1, $zero, _loop`

`lw r2, data` (**hit**)

`call println on r2`

Since `r1 == true`, that validates the not-taken prediction for the branch.

Since `r1 == 0`, the `println` outputs `data=0`!

Memory Consistency Models are often very lax

- Initially, `data == 0, flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread

```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```

- A memory consistency model where above ordering is guaranteed is called, Sequential Consistency (SC): Instructions appear to execute sequentially.
- Real models allow many other orderings to allow optimizations:
 - Write buffers that allow multiple stores to be pending and perform out-of-order
 - Instruction queues that allow loads and other instructions to perform out-of-order
 - Compiler optimizations to reschedule stores and loads out-of-order
- Intel, ARM, Java Virtual Machine all have relaxed memory consistency models
- Moral: **never do custom synchronization** unless you know what you are doing!

Memory Consistency Models are often very lax

- Initially, `data == 0`, `flag == false`.

Producer Thread

```
data == 42;  
flag = true;
```

Consumer Thread


```
while(flag == false) { /* wait */ }  
System.out.println("data=" + data);
```

- Regardless of memory consistency model, they all agree on one thing: that values of `data` and `flag` must be made coherent *eventually*.
 - They only disagree on when that eventually is.
- This property is called **cache coherence**.

Implementing Cache Coherence

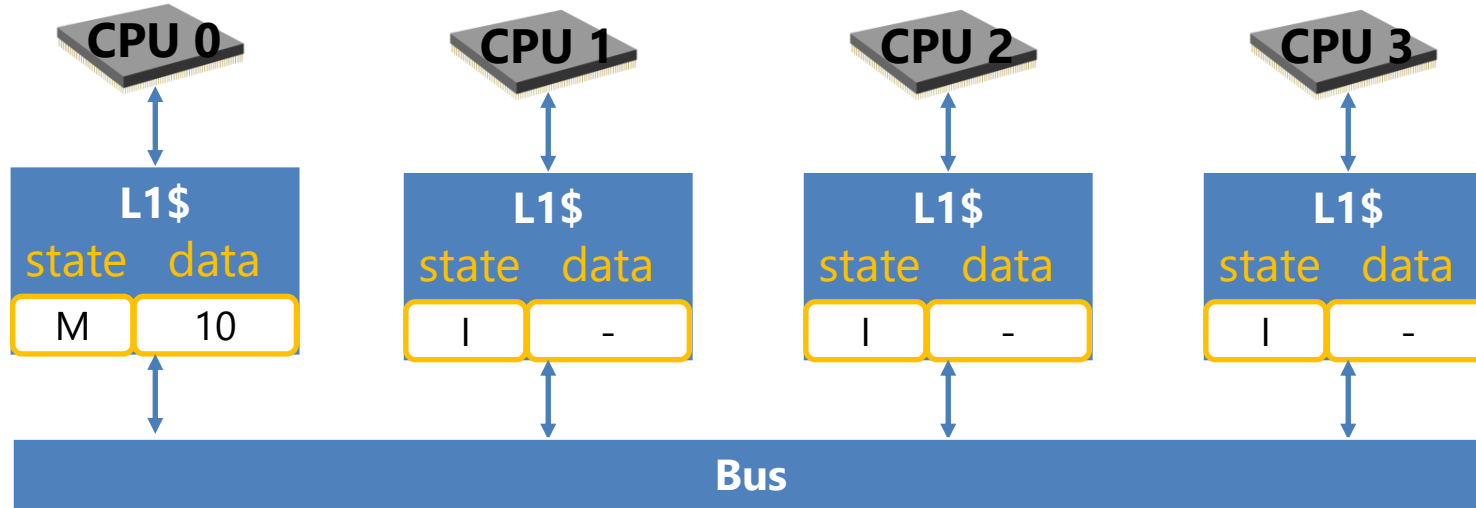
- How to guarantee changes in value are propagated to all caches?
- **Cache coherence protocol**: A protocol, or set of rules, that all caches must follow to ensure coherence between caches
 - MSI (Modified-Shared-Invalid)
 - MESI (Modified-Exclusive-Shared-Invalid)
 - ... often named after the **states** in cache controller **FSM**
- Three states of **MSI** protocol (maintained for each block):
 - Modified: Dirty. Only this cache has copy.
 - Shared: Clean. Other caches may have copy.
 - Invalid: Block contains no data.

MSI Snoopy Cache Coherence Protocol



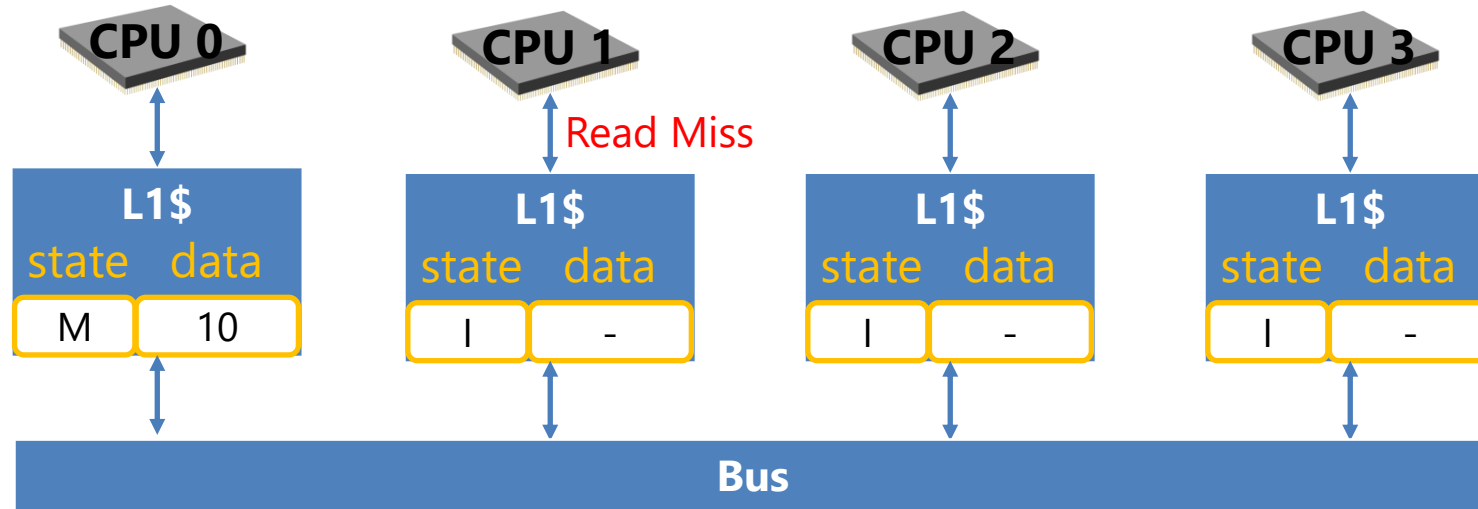
- Each processor **monitors (snoops)** the activity on the **bus**
 - In much the same way as how nodes snoop the Ethernet
- Cache state changes in response to both:
 - Read / writes from the **local processor**
 - Read misses / write misses from **remote processors** it snoops

Scenario 1: Read Snooped on Bus



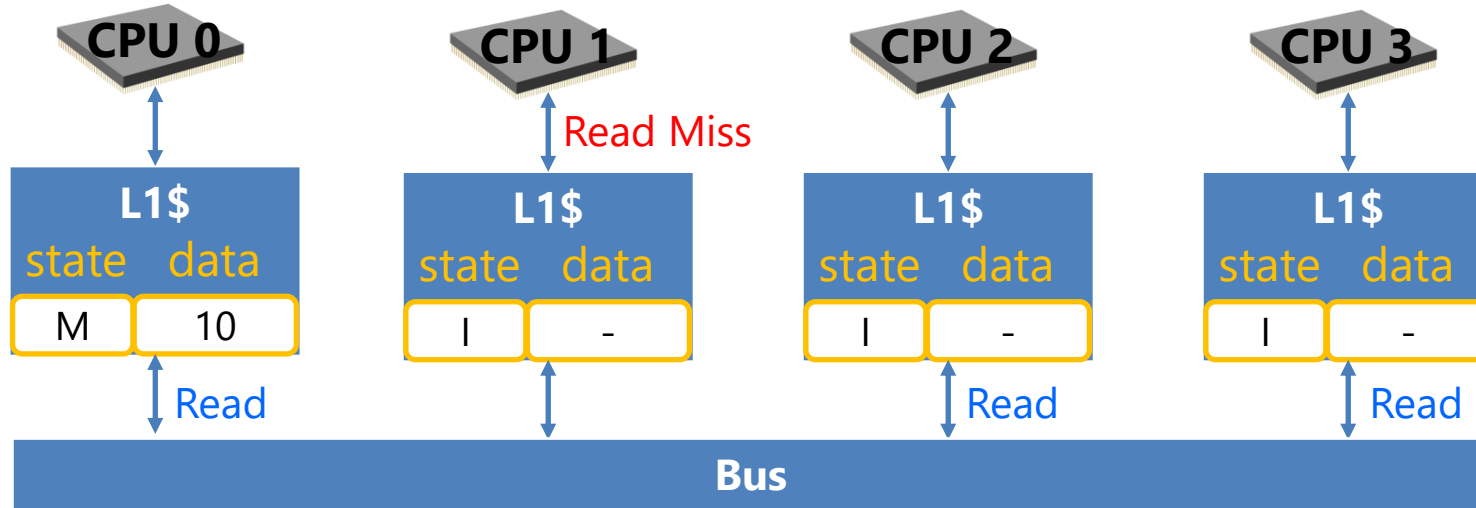
1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state

Scenario 1: Read Snooped on Bus



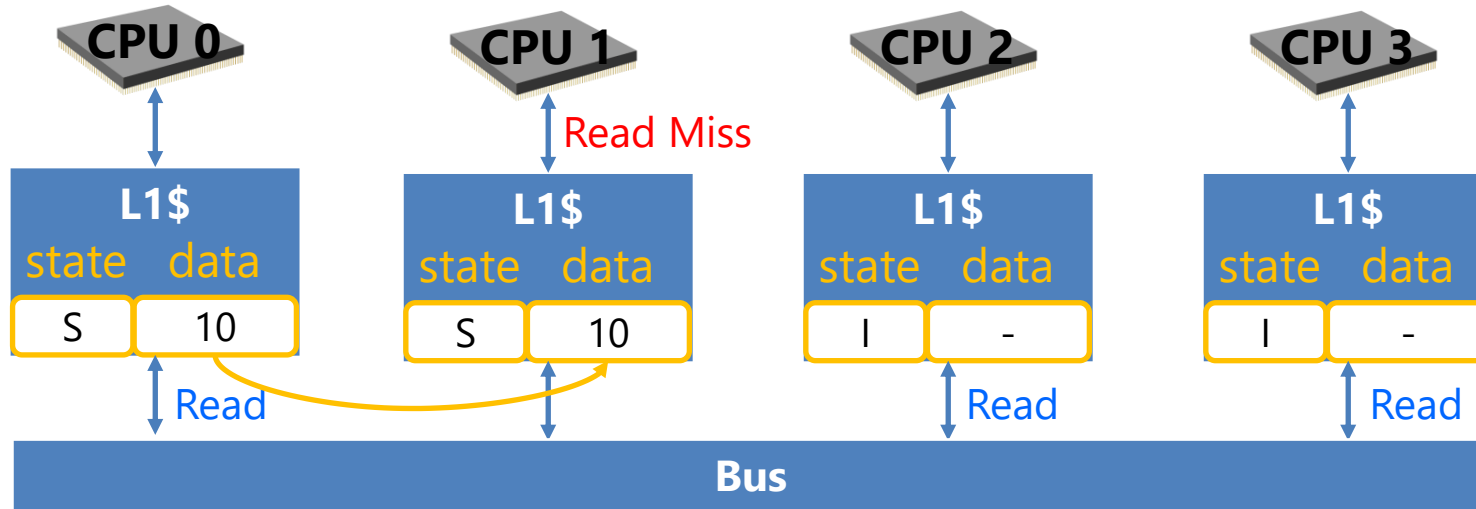
1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a **Read Miss**

Scenario 1: Read Snooped on Bus



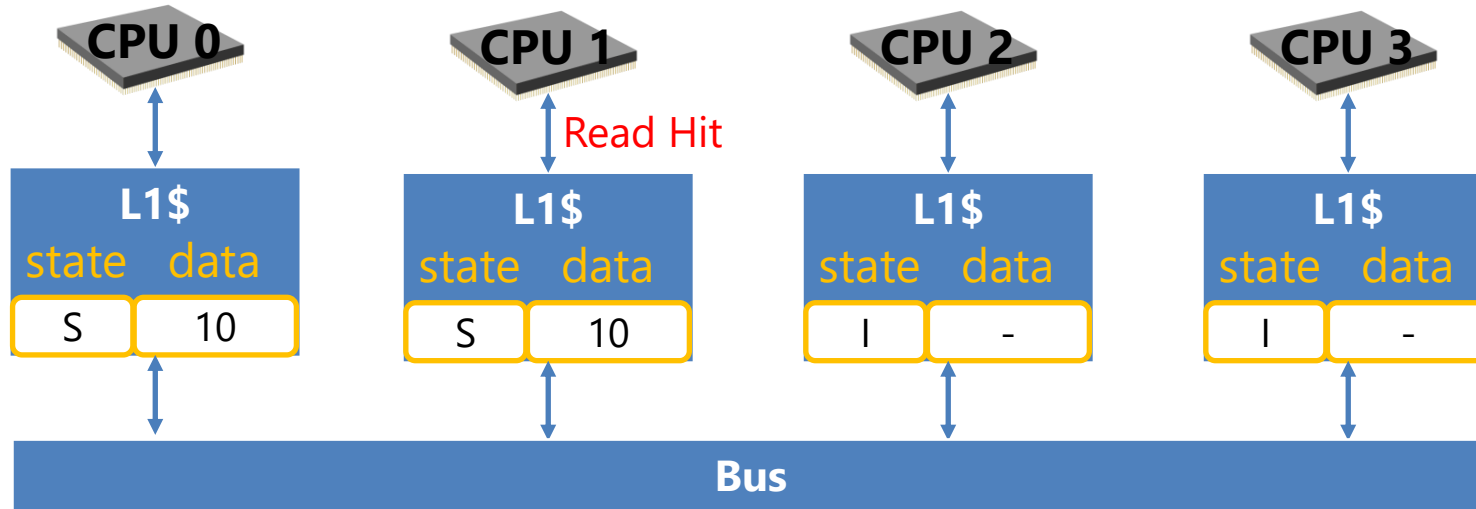
1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a **Read Miss**
3. A **Read** message is broadcast on the Bus

Scenario 1: Read Snooped on Bus



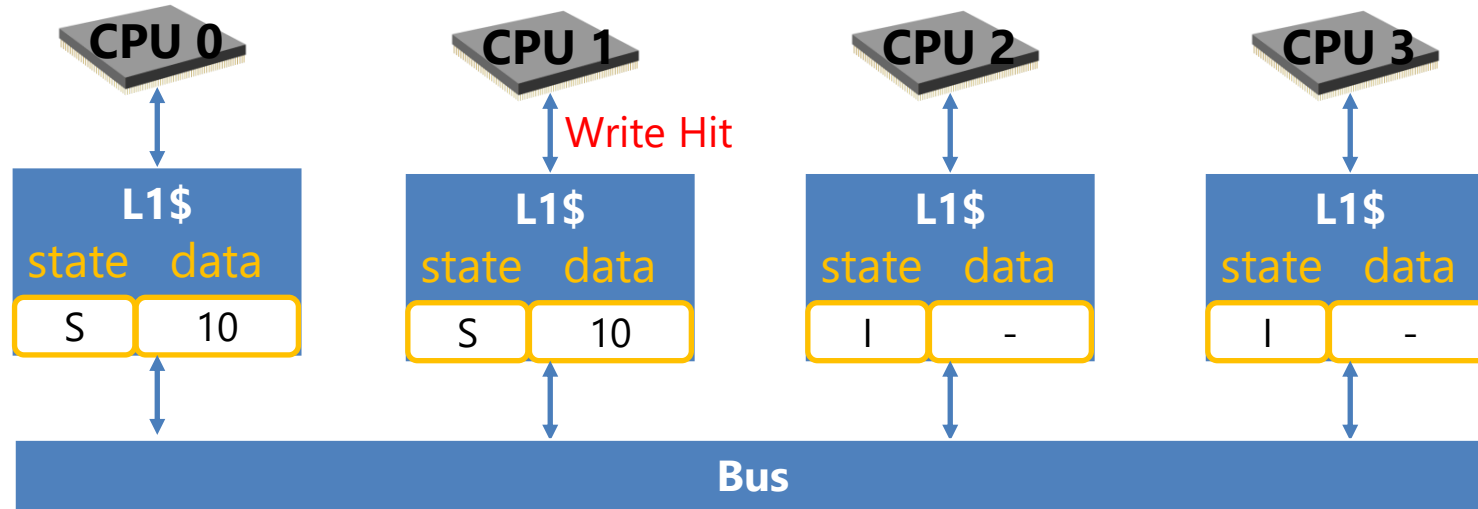
1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a **Read Miss**
3. A **Read** message is broadcast on the Bus
4. CPU 0 snoops message and provides line to CPU1
 - Line in CPU 0 and CPU 1 both transition to Shared state

Scenario 1: Read Snooped on Bus



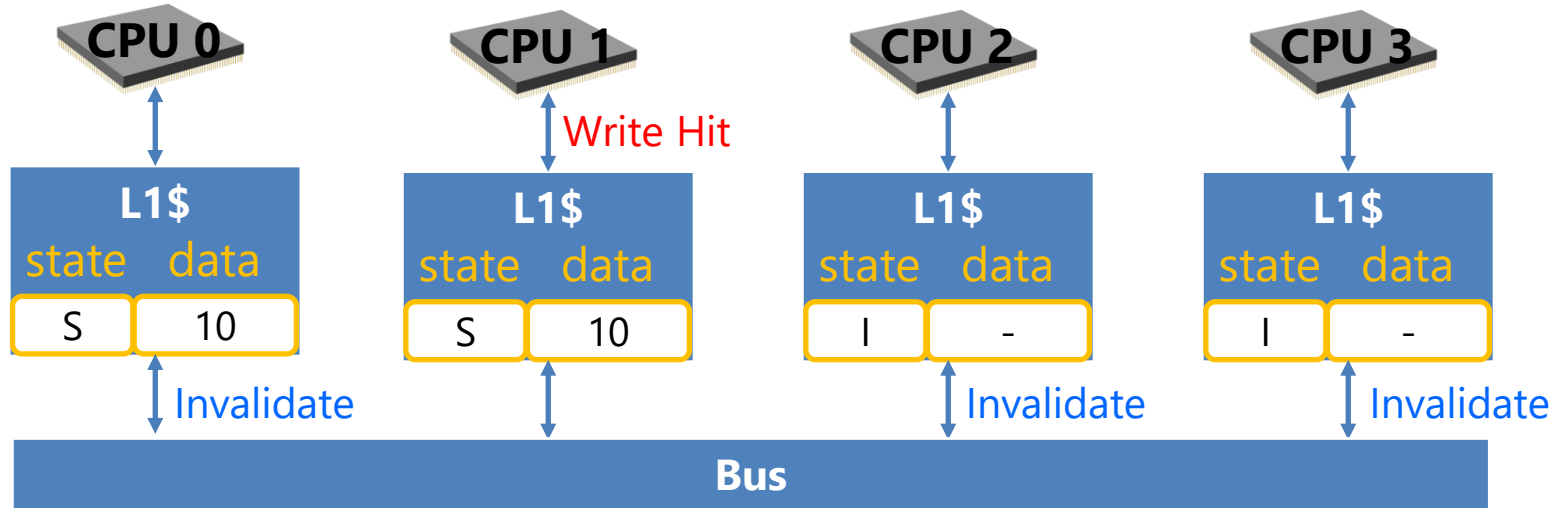
- Subsequent read hits on CPU 0 or CPU 1 don't generate messages
- Only read misses generate messages, not read hits
→ Reduces bus bandwidth pressure since misses are rare!

Scenario 2: Invalidate Snooped on Bus



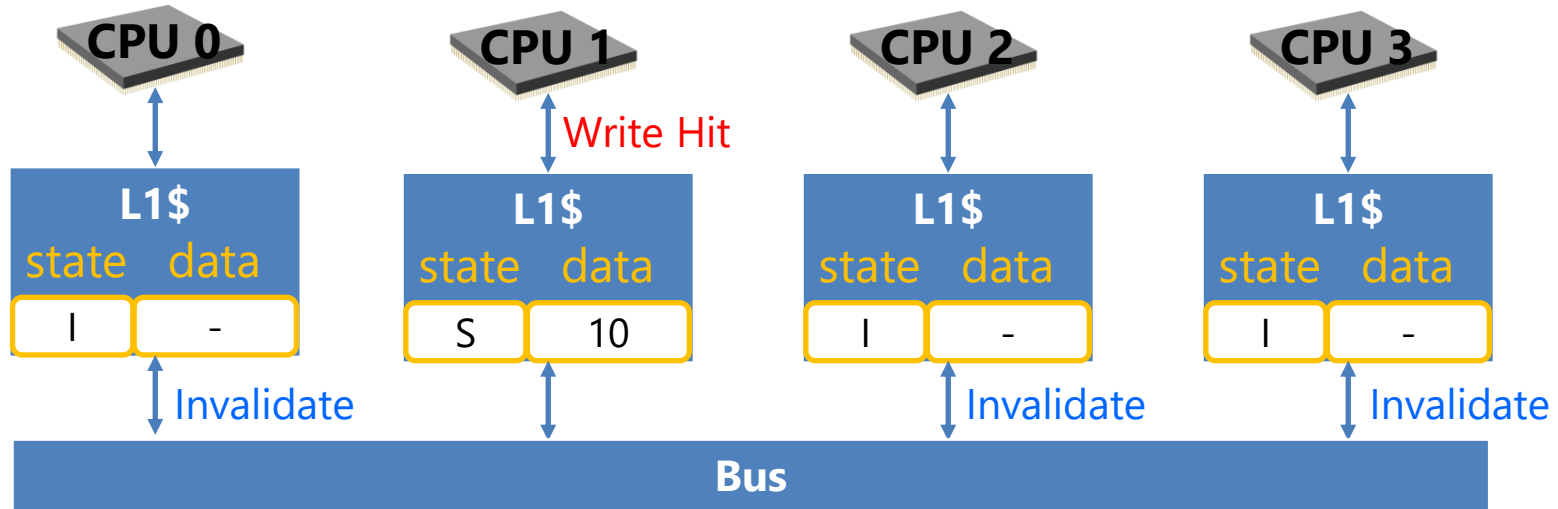
1. CPU 1 writes to line, causing a **Write Hit**

Scenario 2: Invalidate Snooped on Bus



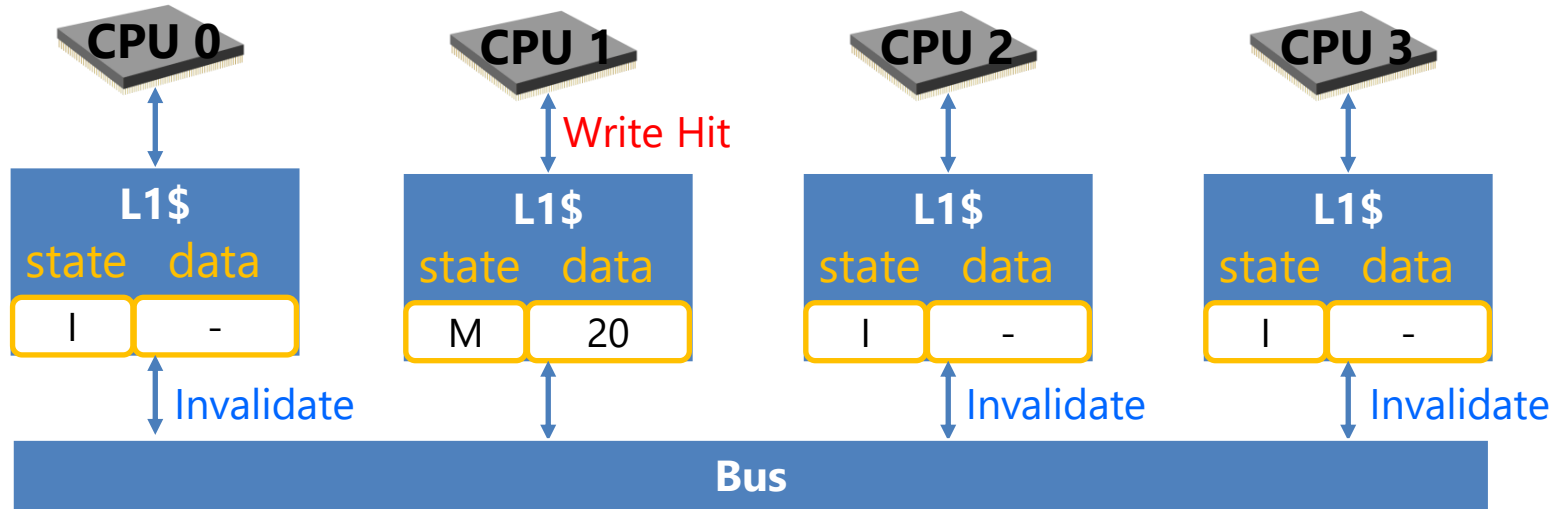
1. CPU 1 writes to line, causing a **Write Hit**
2. An **Invalidate** message is broadcast on the Bus
 - To remove all copies of the line that may become inconsistent

Scenario 2: Invalidate Snooped on Bus



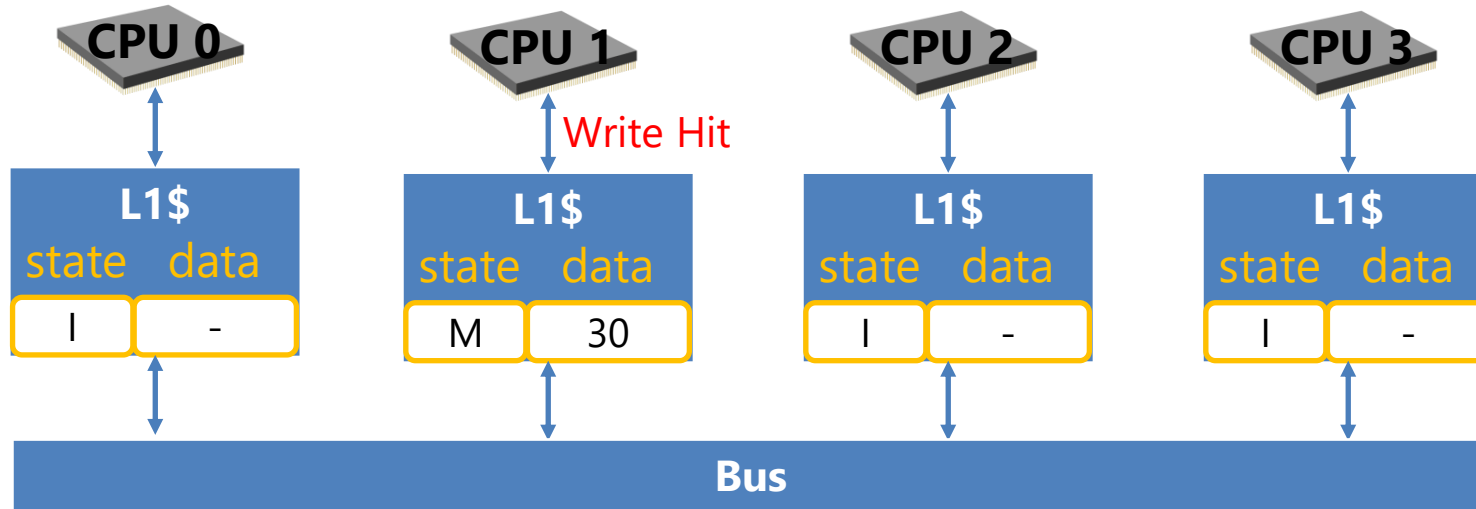
1. CPU 1 writes to line, causing a **Write Hit**
2. An **Invalidate** message is broadcast on the Bus
 - To remove all copies of the line that may become inconsistent
3. CPU 0 snoops message and invalidates its line

Scenario 2: Invalidate Snooped on Bus



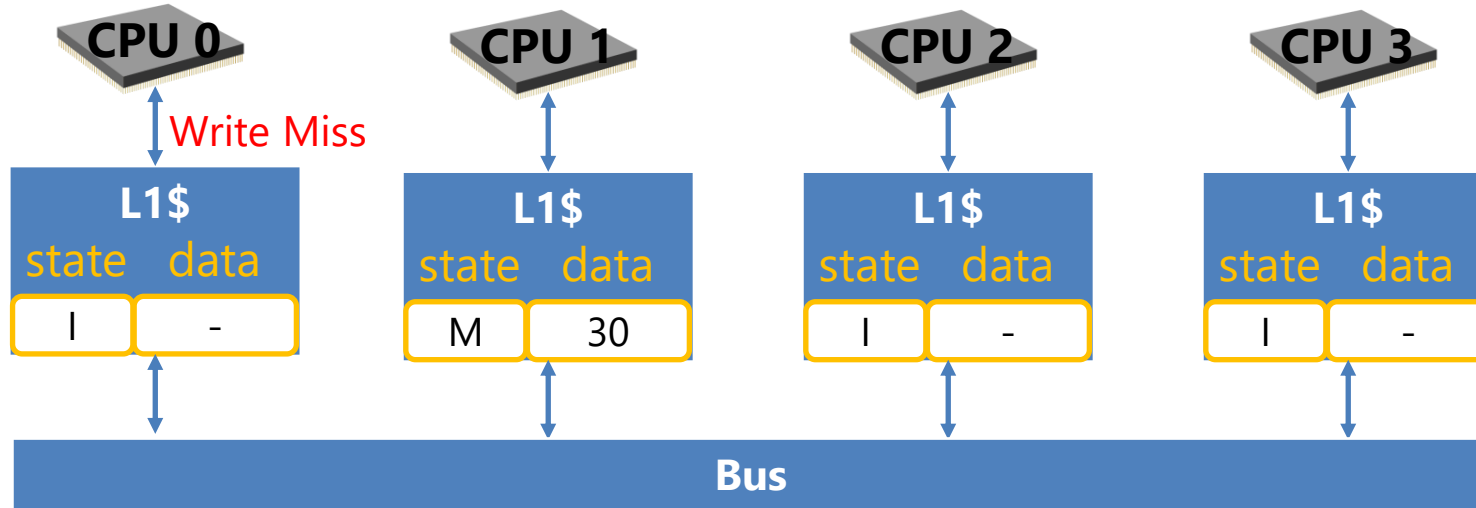
1. CPU 1 writes to line, causing a **Write Hit**
2. An **Invalidate** message is broadcast on the Bus
 - To remove all copies of the line that may become inconsistent
3. CPU 0 snoops message and invalidates its line
4. Line in CPU 1 transitions to Modified state

Scenario 2: Invalidate Snooped on Bus



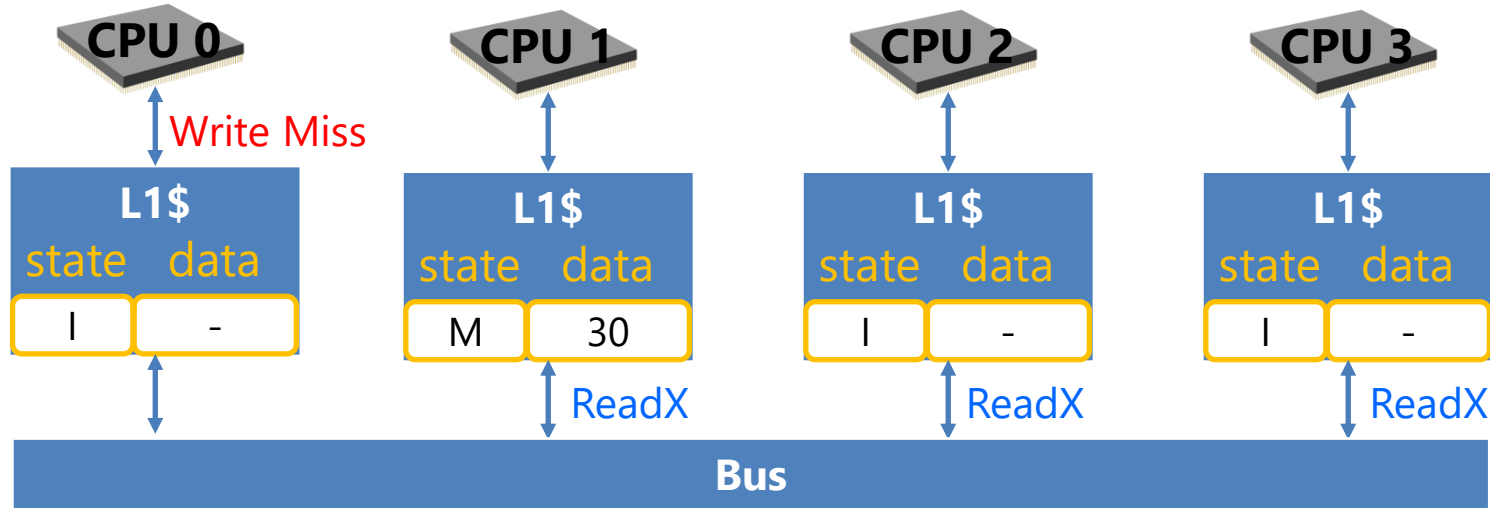
- Subsequent write hits on CPU 1 don't generate messages
- Only write hit in Shared state generate messages
→ Reduces bandwidth since most write hits are to Modified state

Scenario 3: Read Exclusive Snooped on Bus



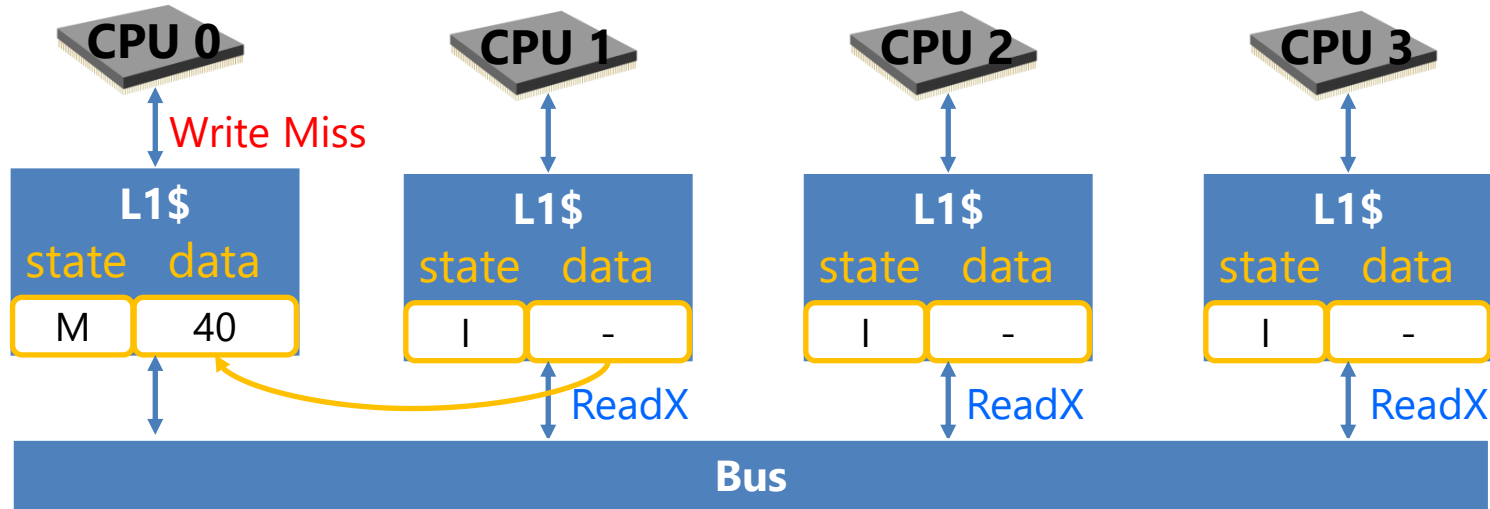
1. CPU 0 writes to line, causing a **Write Miss**

Scenario 3: Read Exclusive Snooped on Bus



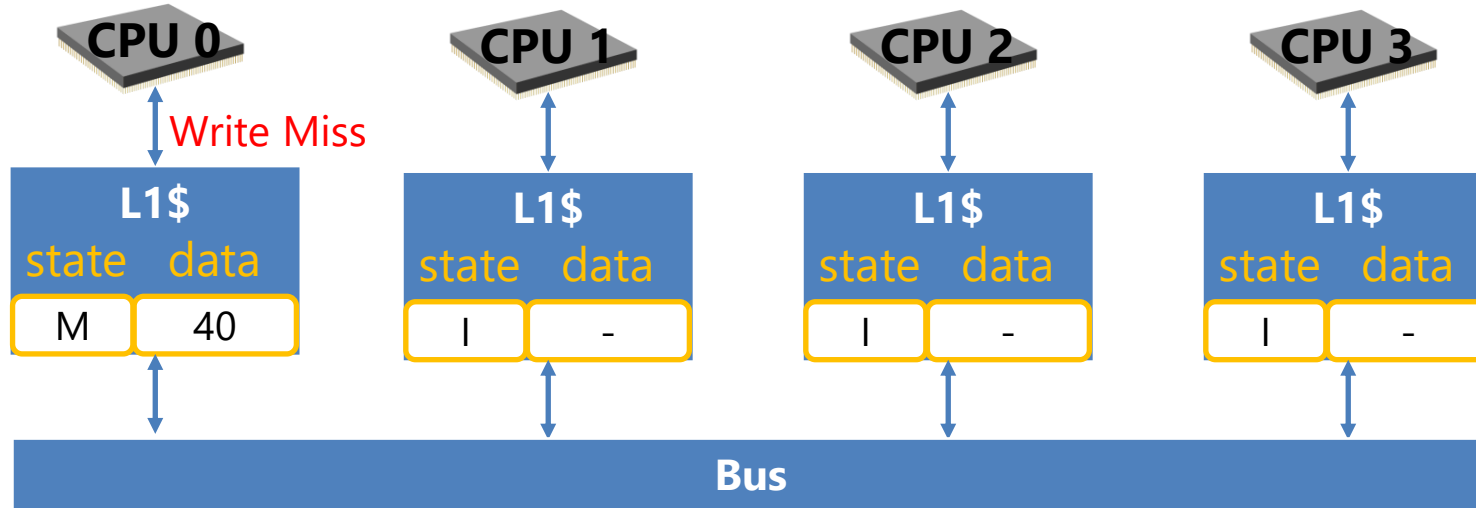
1. CPU 0 writes to line, causing a **Write Miss**
2. A **ReadX** (read exclusive) message is broadcast on the Bus
 - To read line into CPU 0 before writing to it (remember?)
 - To remove all copies of the line that may become inconsistent

Scenario 3: Read Exclusive Snooped on Bus



1. CPU 0 writes to line, causing a **Write Miss**
2. A **ReadX** (read exclusive) message is broadcast on the Bus
 - To read line into CPU 0 before writing to it (remember?)
 - To remove all copies of the line that may become inconsistent
3. CPU 1 snoops message and provides line to CPU 0
 - Line in CPU 1 is invalidated before line in CPU 0 is modified

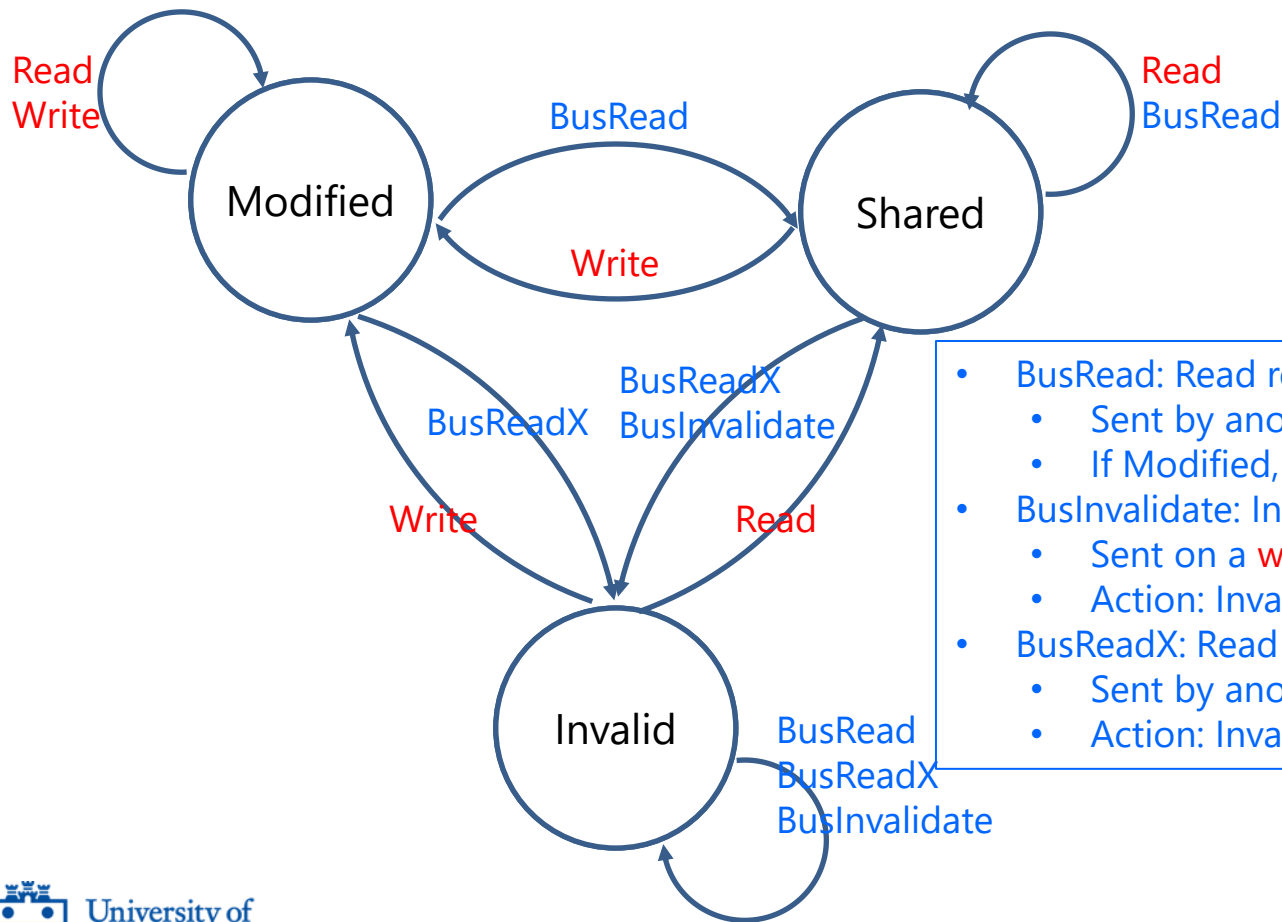
Scenario 3: Read Exclusive Snooped on Bus



- Subsequent write hits on CPU 0 don't generate messages
- Write miss generates message, but not subsequent write hits
→ Reduces bus bandwidth pressure since misses are rare!

Cache Controller FSM for MSI Protocol

- Processor activity in red, Bus activity in blue



- BusRead: Read request is snooped
 - Sent by another processor on a **read miss**
 - If Modified, transition to Shared state
- BusInvalidate: Invalidate request is snooped
 - Sent on a **write hit on shared cache line**
 - Action: Invalidate shared line (if it exists)
- BusReadX: Read exclusive request is snooped
 - Sent by another processor on a **write miss**
 - Action: Invalidate shared line (if it exists)

TLB Coherence

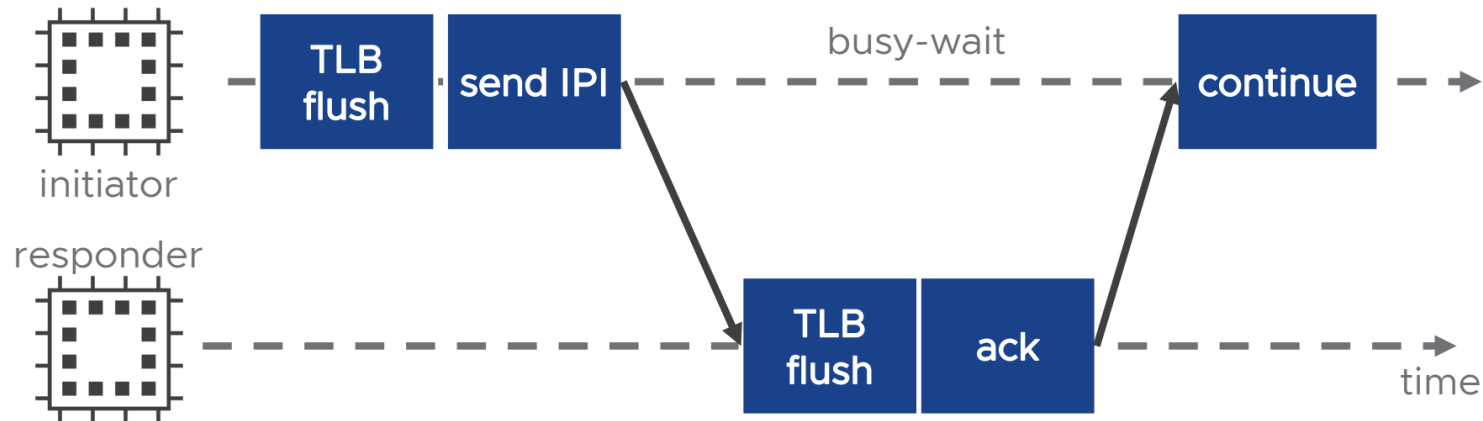
How about TLBs?

- We said TLBs are also a type of cache that caches PTEs.
 - So what happens if a processor changes a PTE?
 - How does that change get propagated to other processor TLBs?
- Unfortunately, there is no hardware coherence for TLBs. ☹️
- That means software (the OS) must handle the coherence
 - Which is of course much much slower

TLB shutdown

- In order to update a PTE (page table entry)
 - Initiator OS must first flush its own TLB
 - Send IPIs (Inter-processor interrupts) to other processors
 - To flush the TLBs for all other processors too
 - Source of significant performance overhead

TLB Flushes in Linux and FreeBSD



* Courtesy of Nadav Amit et al. at VMWare