Multiprocessors and Caching

CS 1541
Wonsun Ahn

University of Pittsburgh

Two ways to use multiple processors

- **Distributed (Memory) System**
  - Processors **do not share memory** (and by extension **data**)
  - Processors exchange data through network messages
  - Programming standards:
    - Message Passing Interface (MPI) – C/C++ API for exchanging messages
    - Ajax (Asynchronous JavaScript and XML) – API for web apps
  - Data exchange protocols: TCP/IP, UDP/IP, JSON, XML…

- **Shared Memory System** (a.k.a. Multiprocessor System)
  - Processors **share memory** (and by extension **data**)
  - Programming standards:
    - Pthreads (POSIX threads), Java threads – APIs for threading
    - OpenMP – Compiler #pragma directives for parallelization
  - **Cache coherence protocol**: protocol for exchanging data among caches

→ Just like Ethernet, caches are part of a larger network of caches

Shared Data Review

- What bad thing can happen when you have shared data?

- Dataraces!
  - You should have learned it in CS 449.
  - But if you didn't, don't worry I'll go over it.

Review: Datarace Example

```
int shared = 0;
void *add(void *unused) {
  for(int i=0; i < 1000000; i++) { shared++; }
  return NULL;
}
int main() {
  pthread_t t;
  // Child thread starts running add
  pthread_create(&t, NULL, add, NULL);
  // Main thread starts running add
  add(NULL);
  // Wait until child thread completes
  pthread_join(t, NULL);
  printf("shared=%d\n", shared);
  return 0;
}
```

```
bash-4.2$ ./datarace
shared=1085894
bash-4.2$ ./datarace
shared=1101173
bash-4.2$ ./datarace
shared=1065494
```

Q) What do you expect from running this? Maybe shared=2000000 ?

A) Nondeterministic result! Due to **datarace** on `shared`.

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

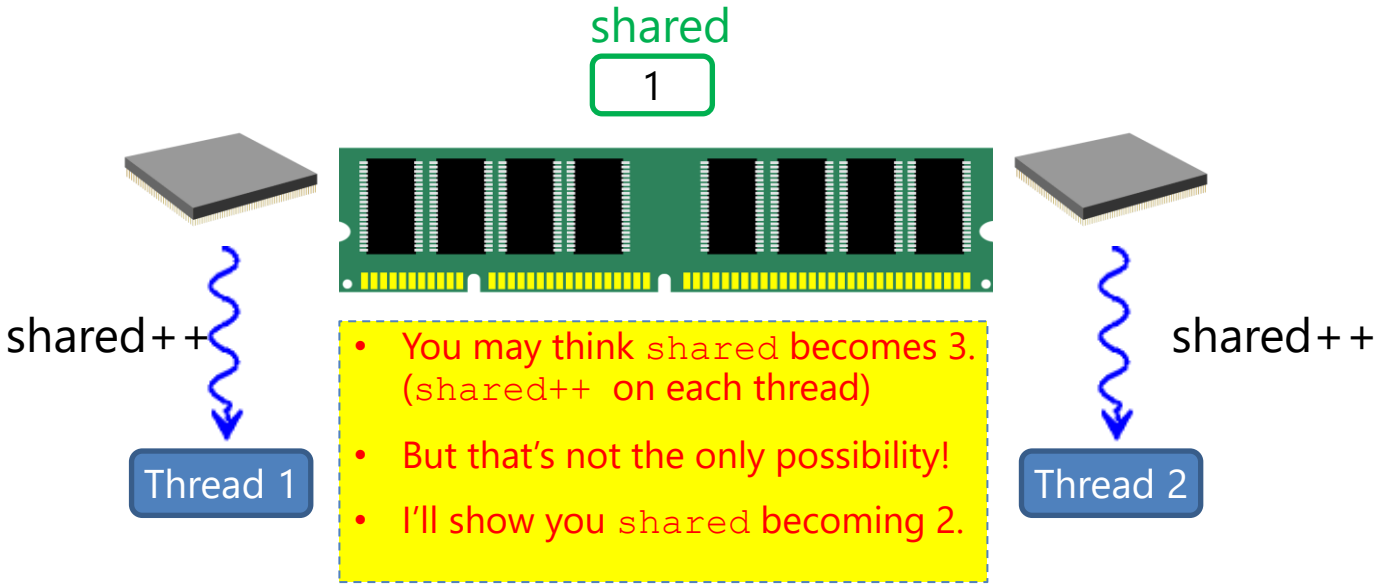

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

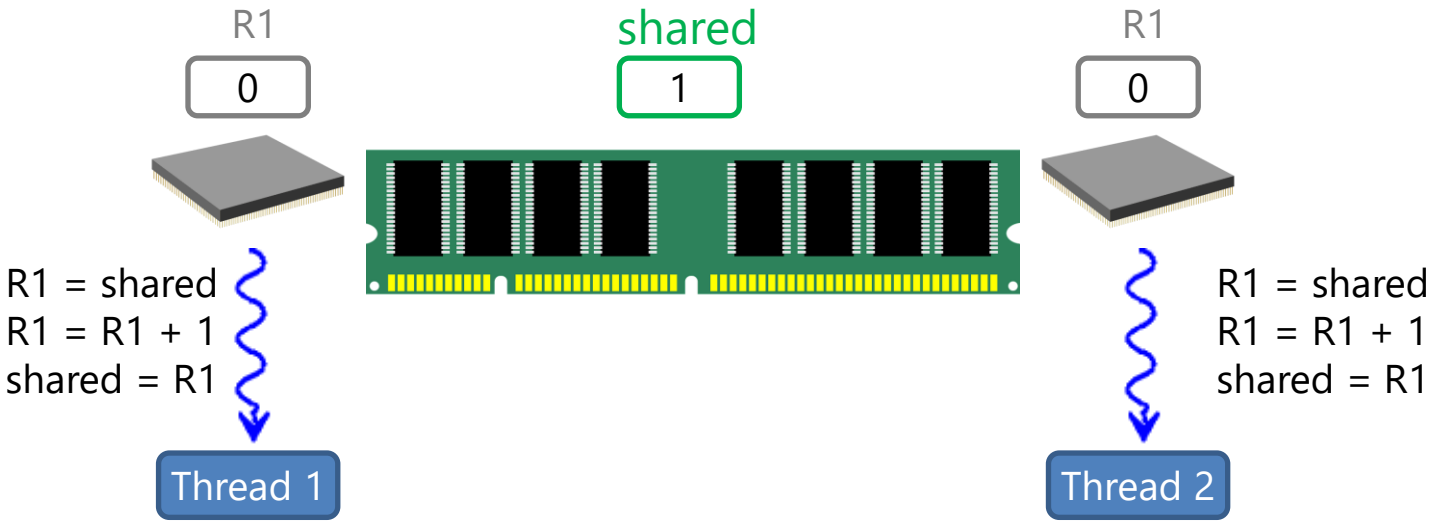

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

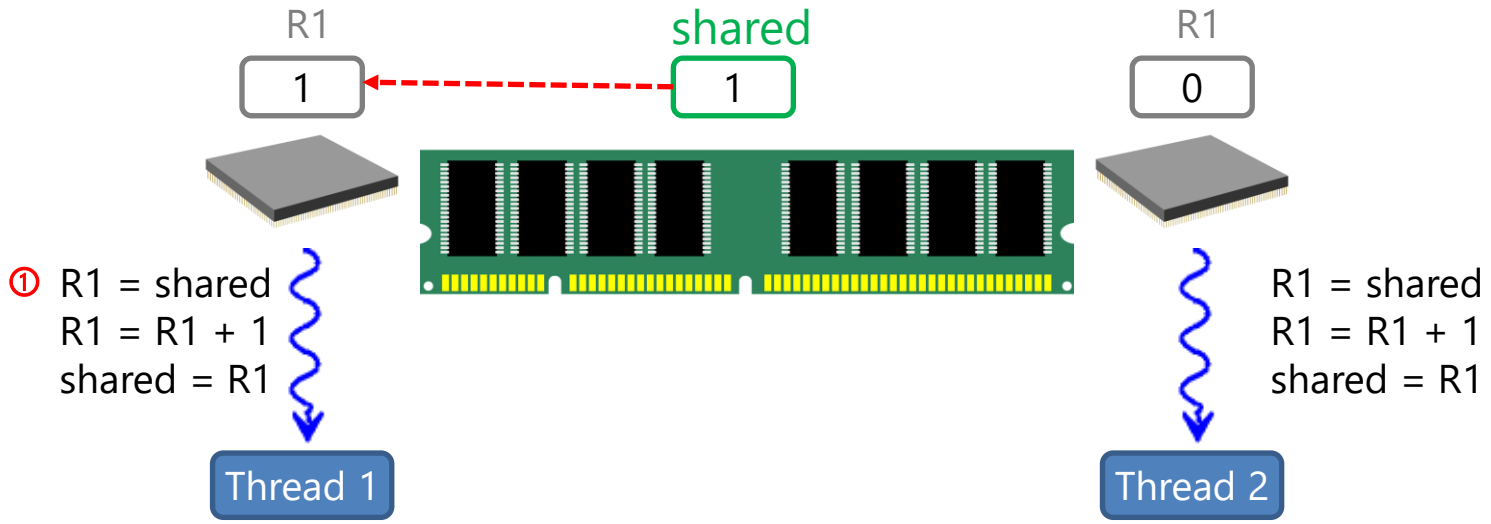

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

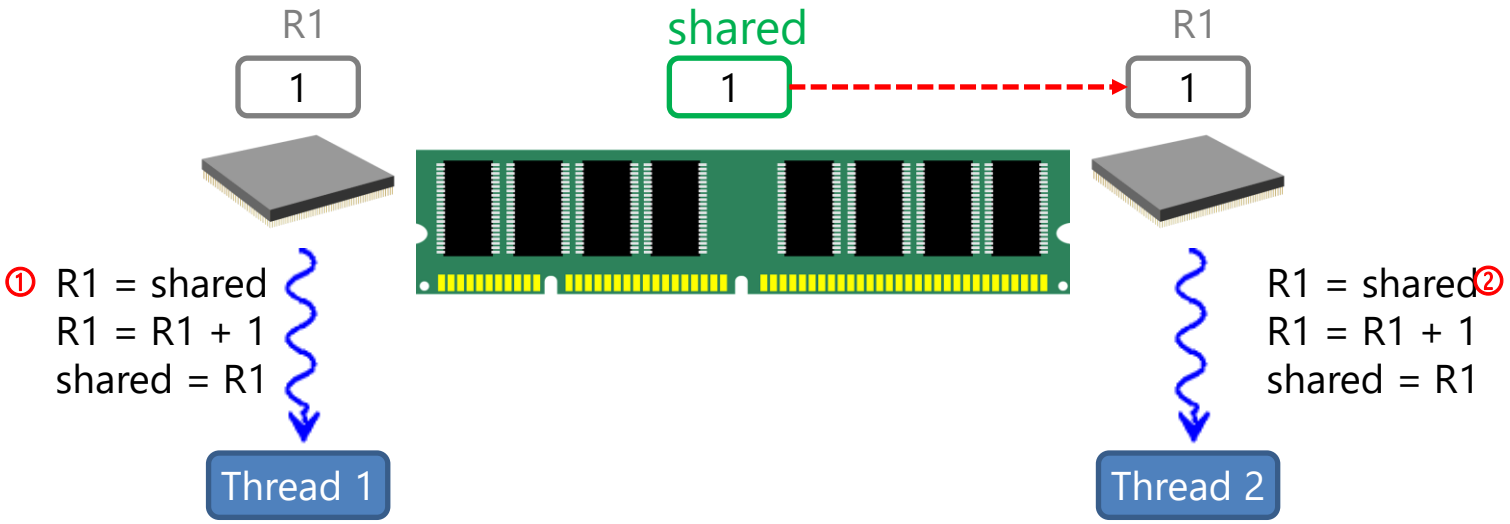

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

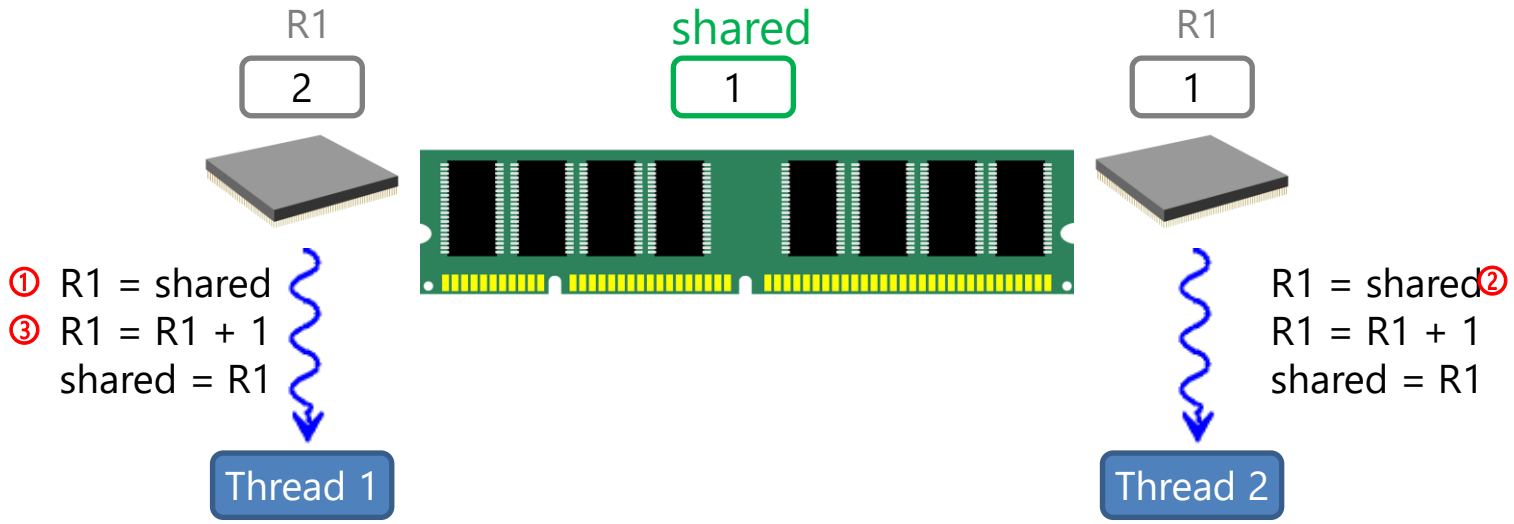

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

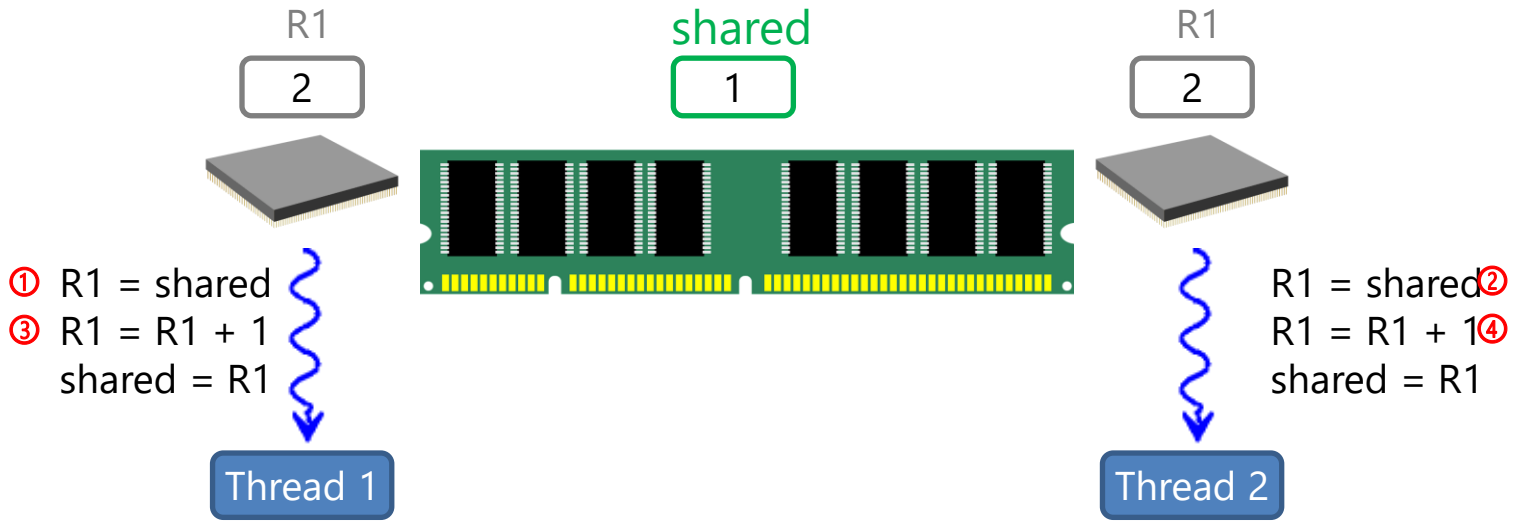

Review: Datarace Example

- When two threads do `shared++;` initially `shared` = 1

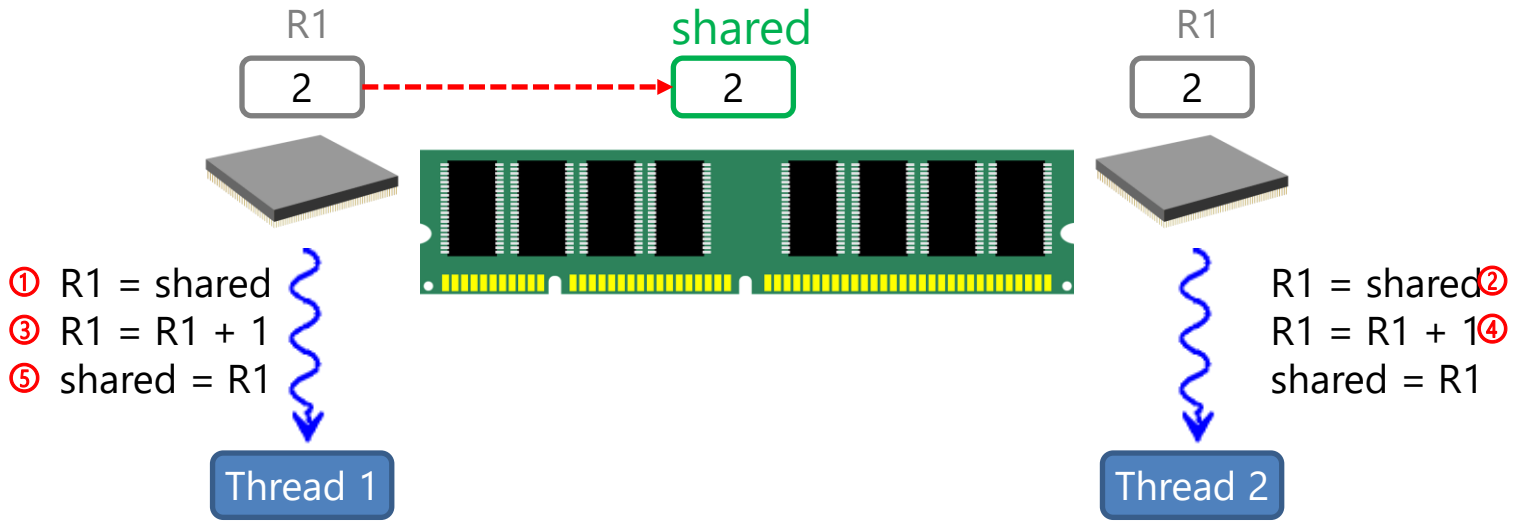

Review: Datarace Example

- Why did this occur in the first place?
- Because data was **replicated** to CPU registers and each worked on its own copy!

- When two threads do `shared++;` initially `shared` = 1

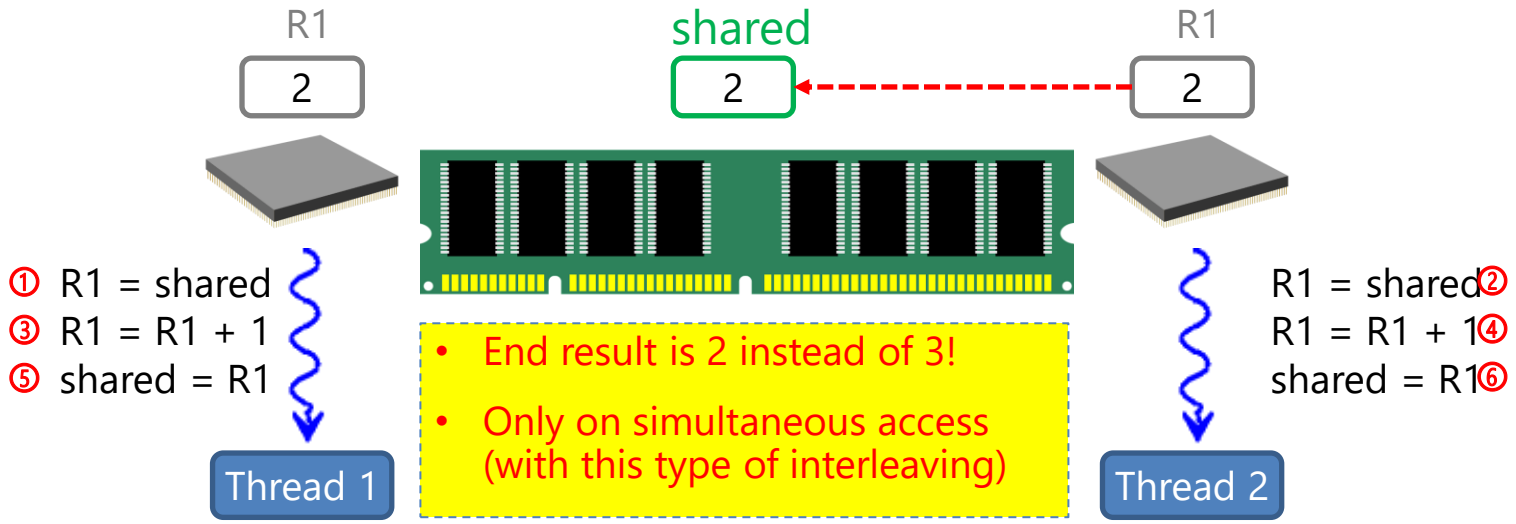

Review: Datarace Example

```
pthread_mutex_t lock;
int shared = 0;
void *add(void *unused) {
  for(int i=0; i < 1000000; i++) {
      pthread_mutex_lock(&lock);
      shared++;
      pthread_mutex_unlock(&lock);
  }
  return NULL;
}
int main() {
  …
}
```

```
bash-4.2$ ./datarace
shared=2000000
bash-4.2$ ./datarace
shared=2000000
bash-4.2$ ./datarace
shared=2000000
```

- Data race is fixed! Now shared is always 2000000.
- Problem solved? No! CPU registers is not the only place **replication** happens!

Caching also does replication!

- What happens if caches sit in between processors and memory?

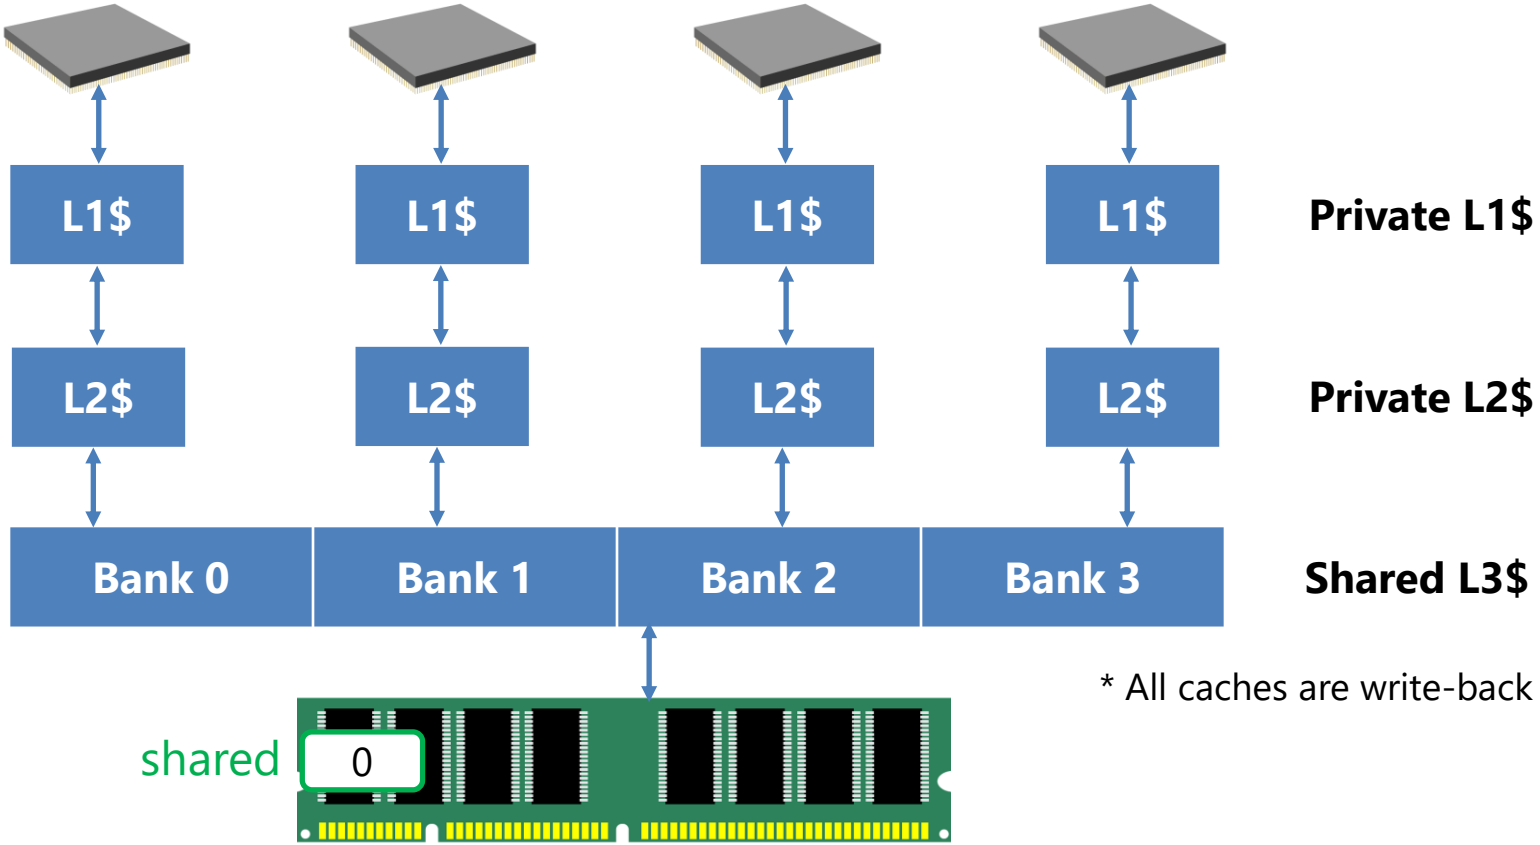

Caching also does replication!

- Let's say CPU 0 first fetches `shared` for incrementing

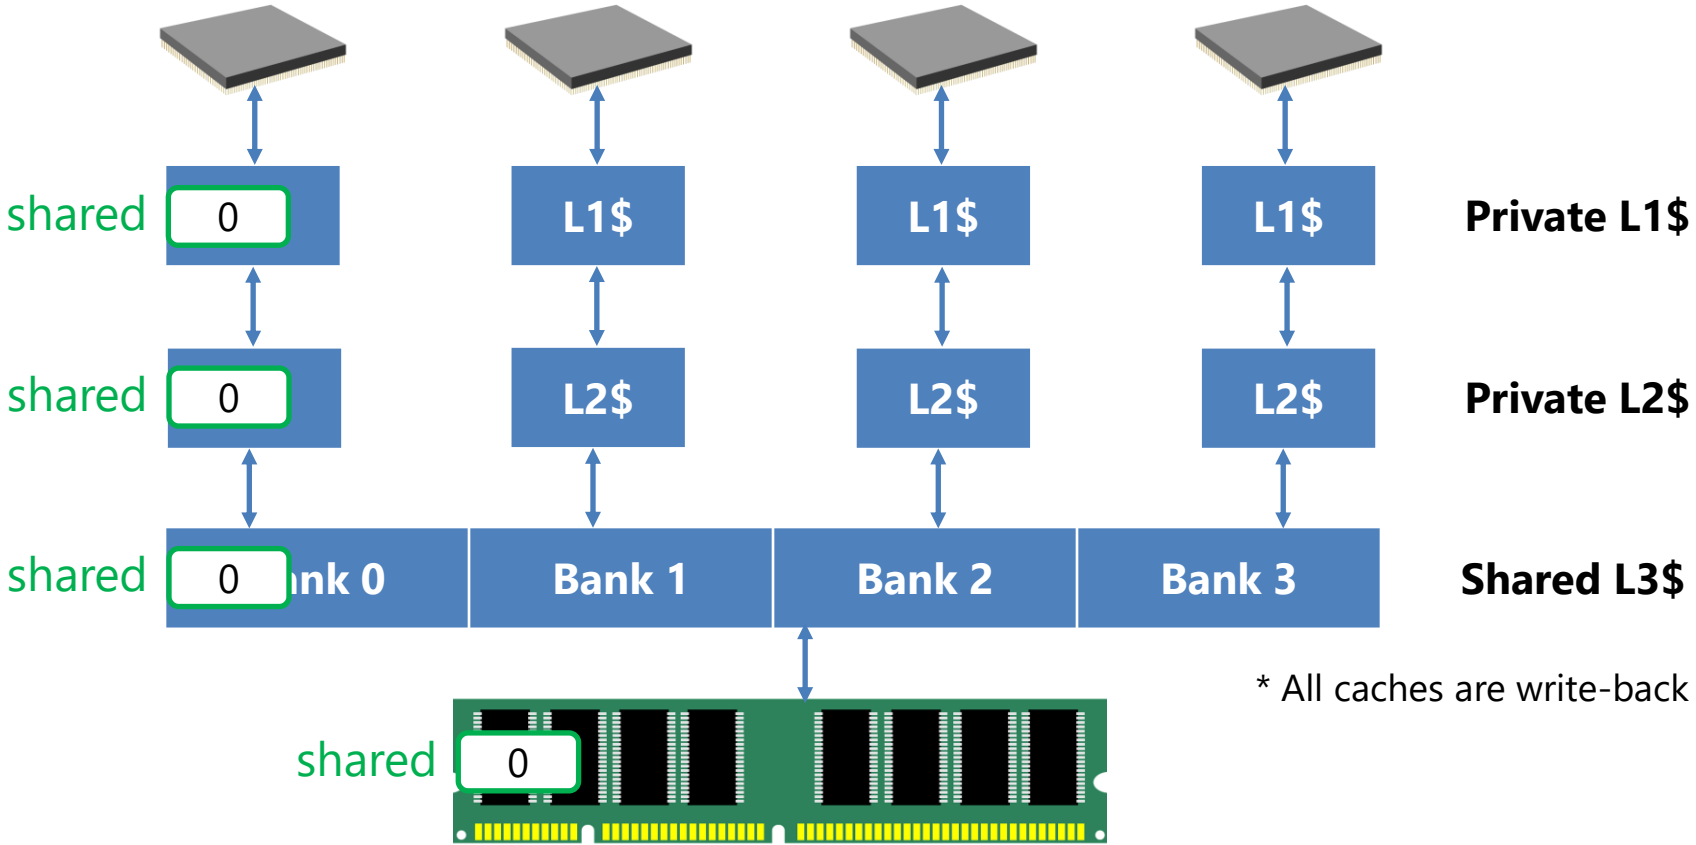

Caching also does replication!

- Then CPU 0 increments `shared` 100 times to 100

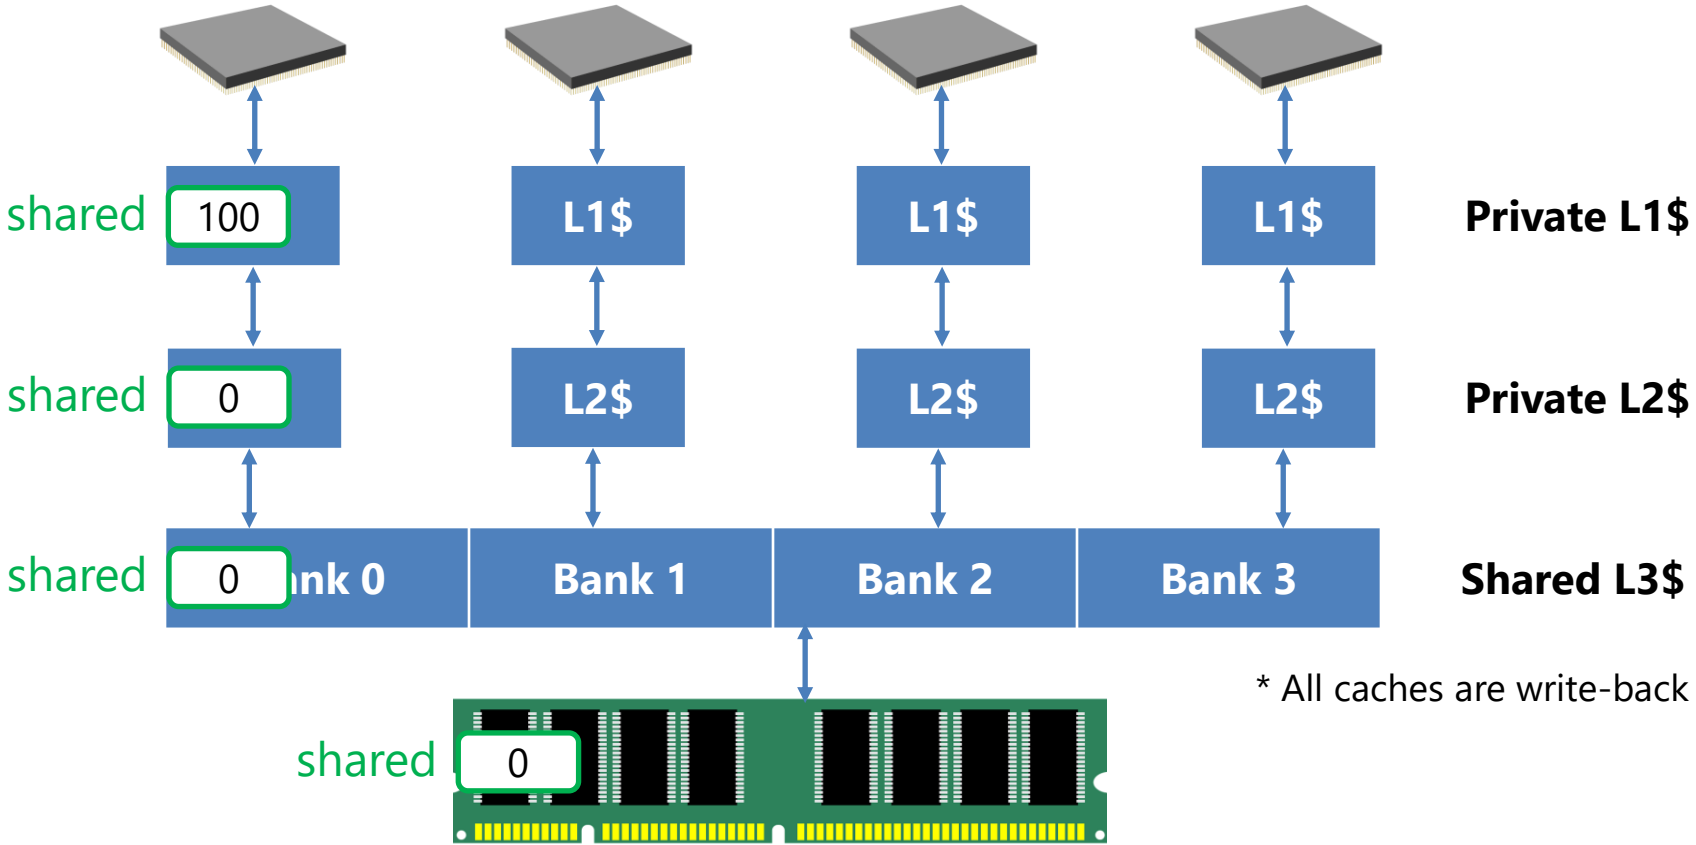

Caching also does replication!

- Then CPU 2 gets hold of the mutex and fetches `shared` from L3

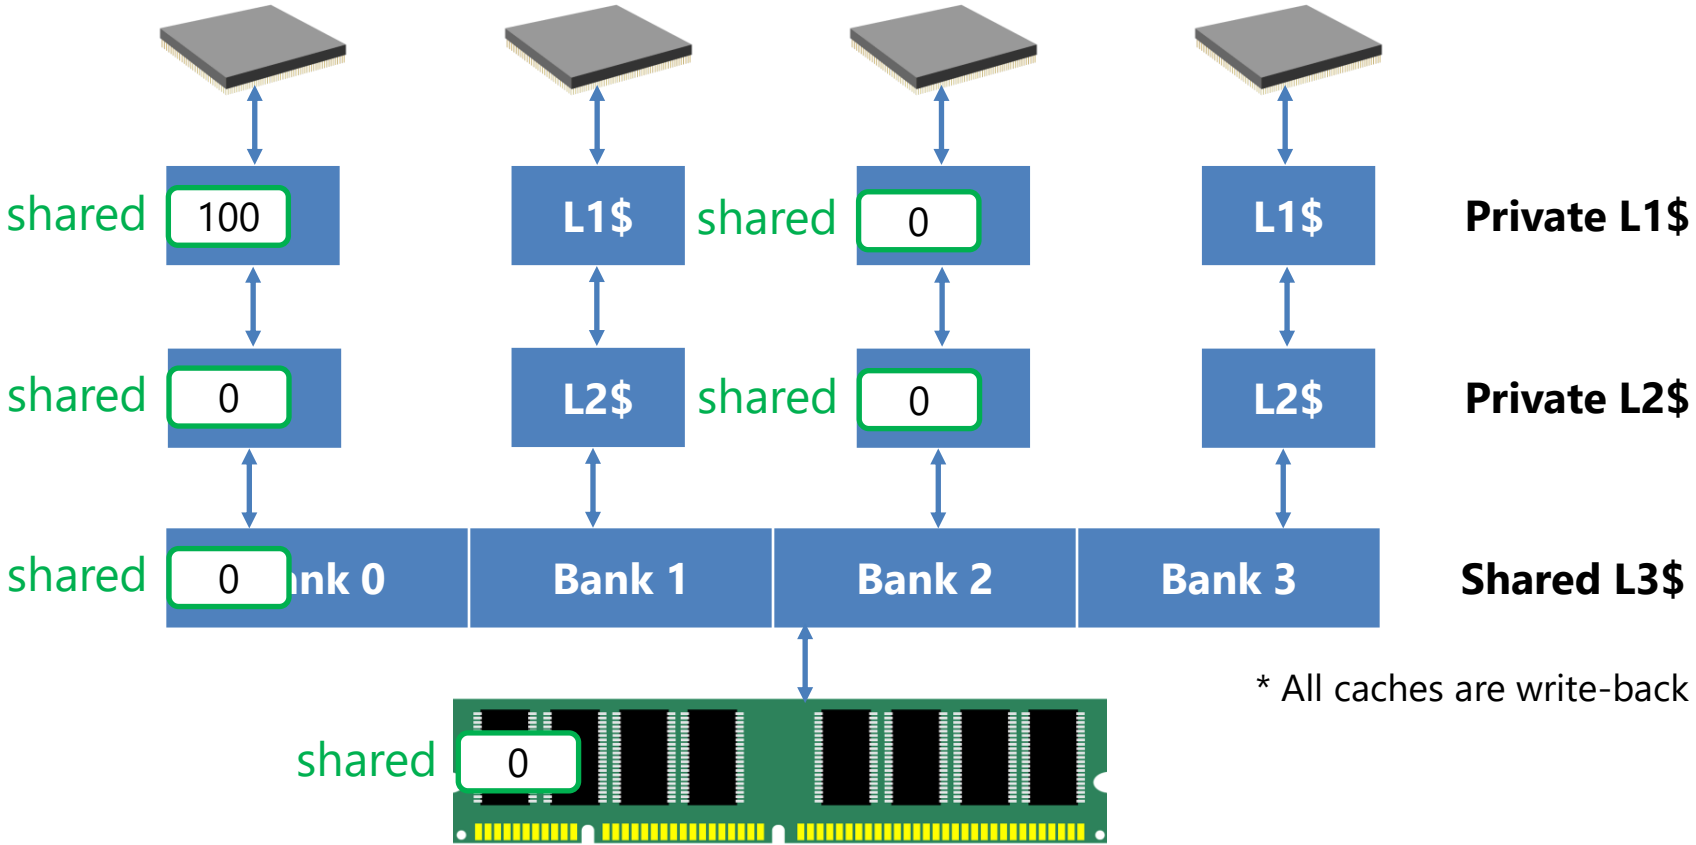

Caching also does replication!

- Then CPU 2 increments `shared` 10 times to 10

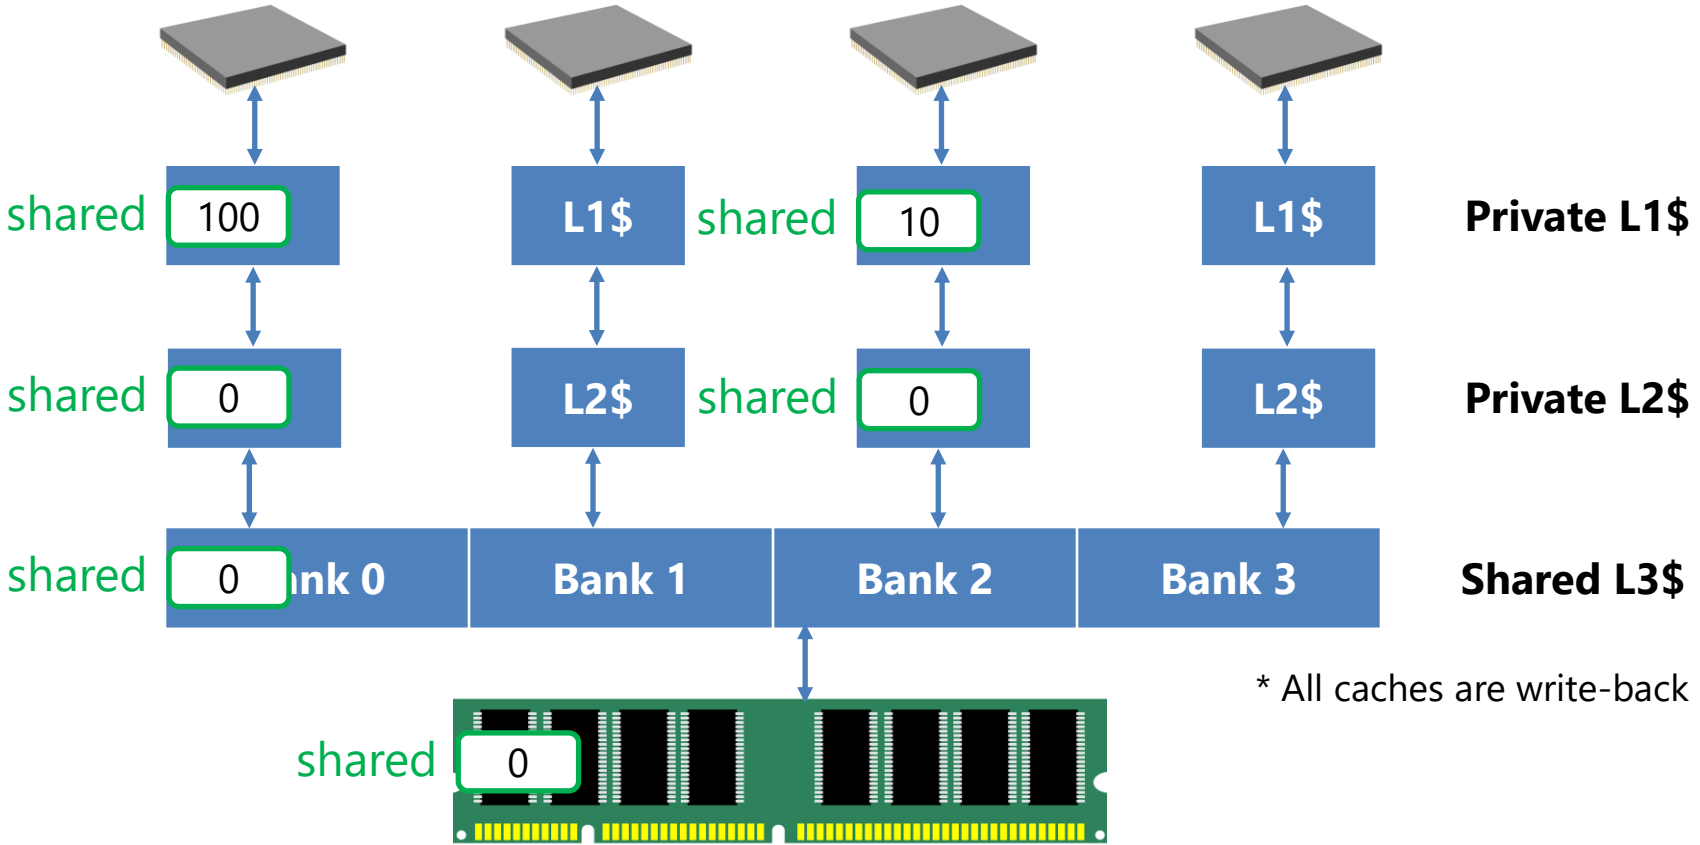

Caching also does replication!

- Clearly this is wrong. L1 caches of CPU 0 and CPU 2 are **incoherent**.

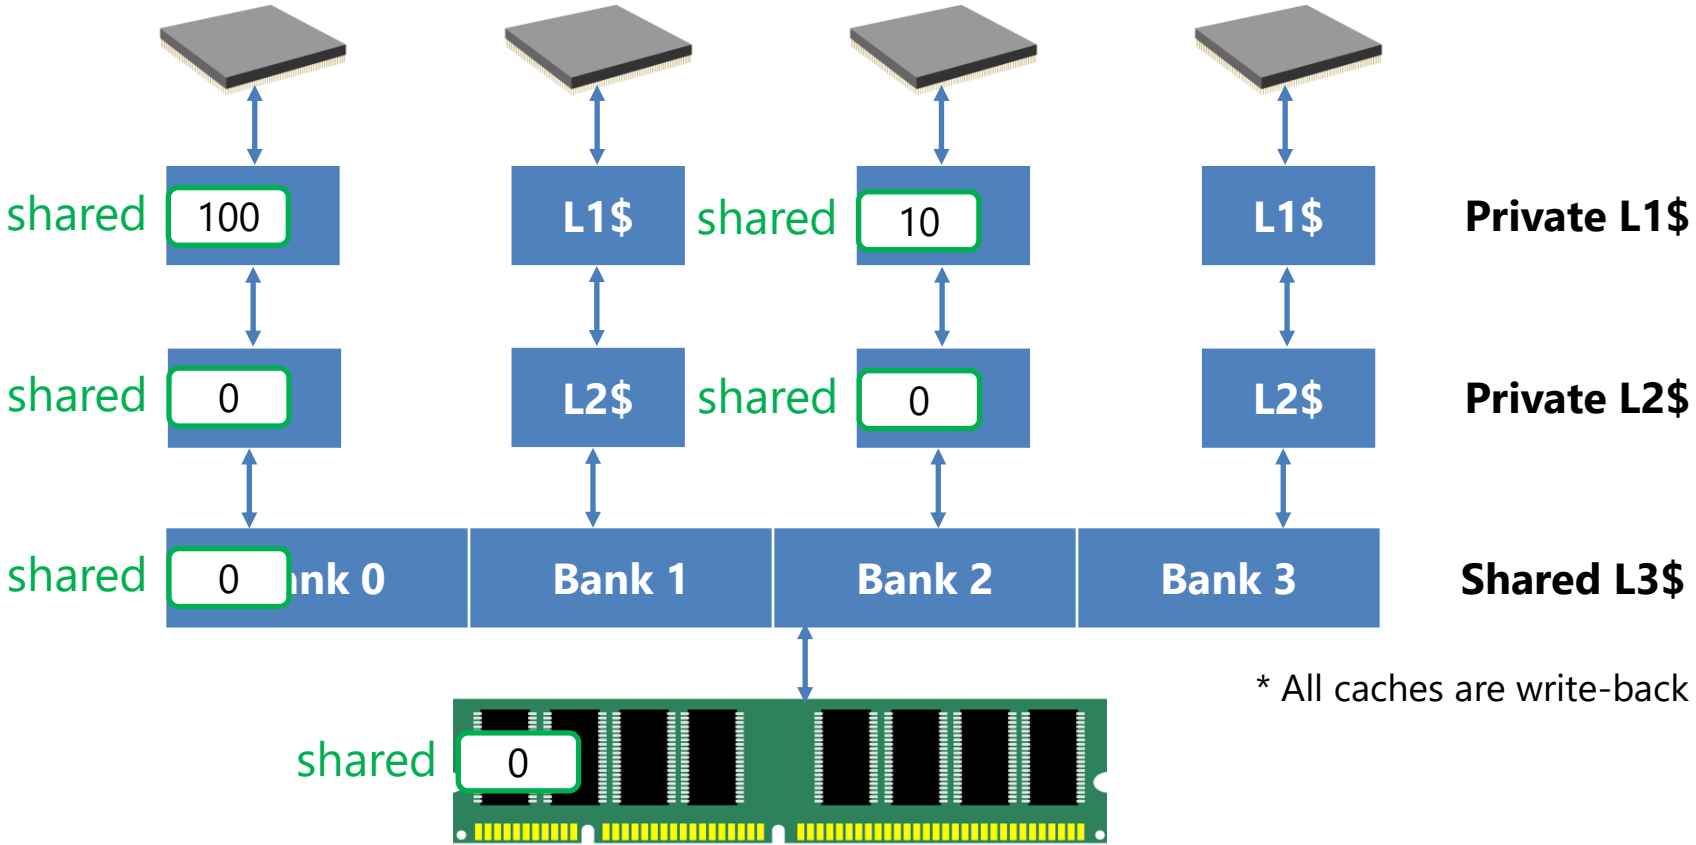

Cache Incoherence: Problem with Private Caches

- This problem does not occur with a **shared cache**.
  - All processors share and work on a **single copy** of data.

- The problem exists for **private** caches.
  - Private copy is at times inconsistent with lower memory.
  - **Incoherence** occurs when private copies differ from each other.
    → Means processors return different values for same location!

Cache Coherence

Cache Coherence

- **Cache coherence** (loosely defined):
  - All processors of system should see the **same view of memory**
  - Copies of values cached by processors should adhere to this rule

- Each ISA has a different definition of what that "view" means
  - **Memory consistency model**: definition of what that "view" is

- All models agree on one thing:
  - That a change in value should reflect on all copies (eventually)

Implementing Cache Coherence

- How to guarantee changes in value are propagated to all caches?

- **Cache coherence protocol**: A protocol, or set of rules, that all caches must follow to ensure coherence between caches
  - MSI (Modified-Shared-Invalid)
  - MESI (Modified-Exclusive-Shared-Invalid)
  - … often named after the **states** in cache controller **FSM**

- Three states of **MSI** protocol (maintained for each block):
  - Modified: Dirty. Only this cache has copy.
  - Shared: Clean. Other caches may have copy.
  - Invalid: Block contains no data.

MSI Snoopy Cache Coherence Protocol

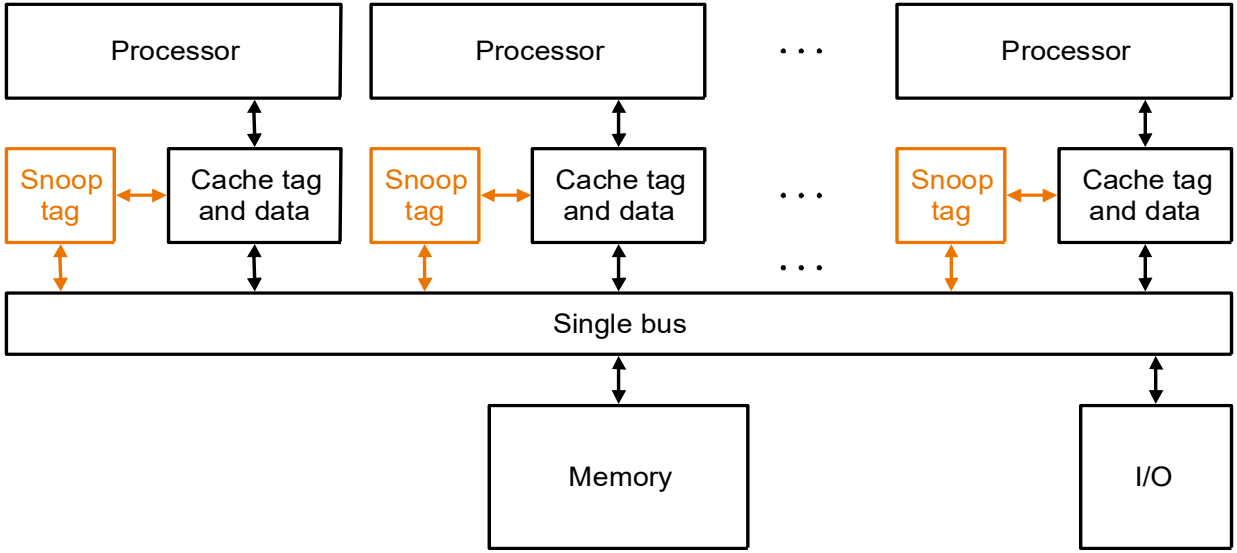

- Each processor **monitors (snoops)** the activity on the **bus**
  - In much the same way as how nodes snoop the Ethernet
- Cache state changes in response to both:
  - Read / writes from the **local processor**
  - Read misses / write misses from **remote processors** it snoops

Scenario 1: Read Snooped on Bus

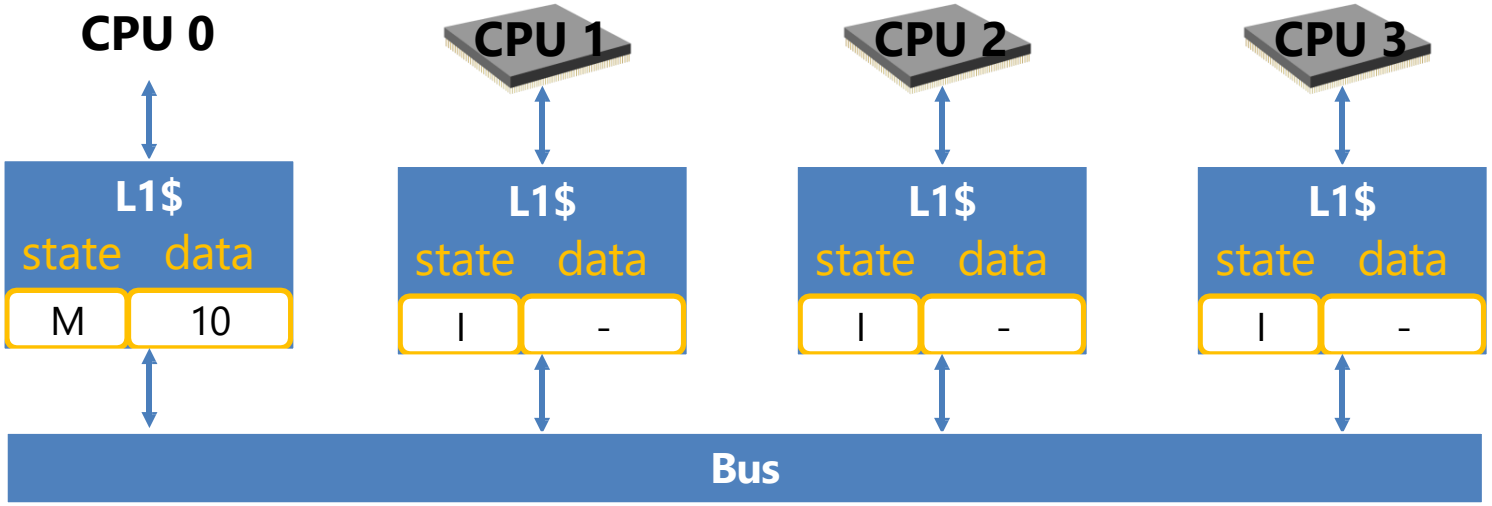

1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state

Scenario 1: Read Snooped on Bus

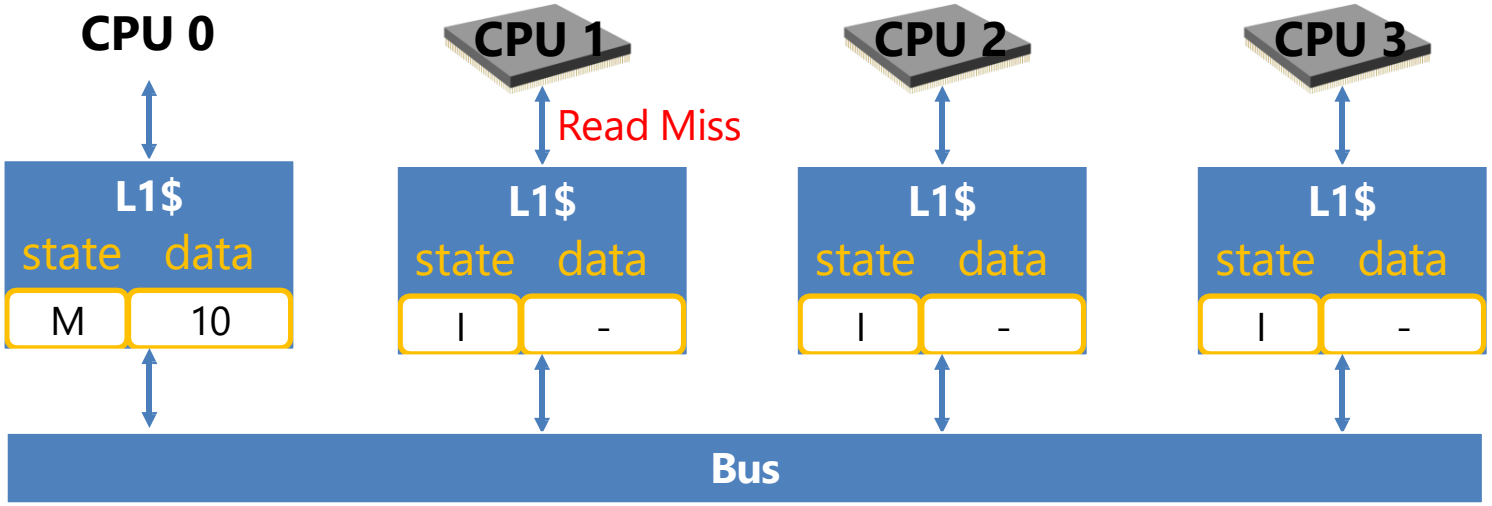

1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a Read Miss

Scenario 1: Read Snooped on Bus

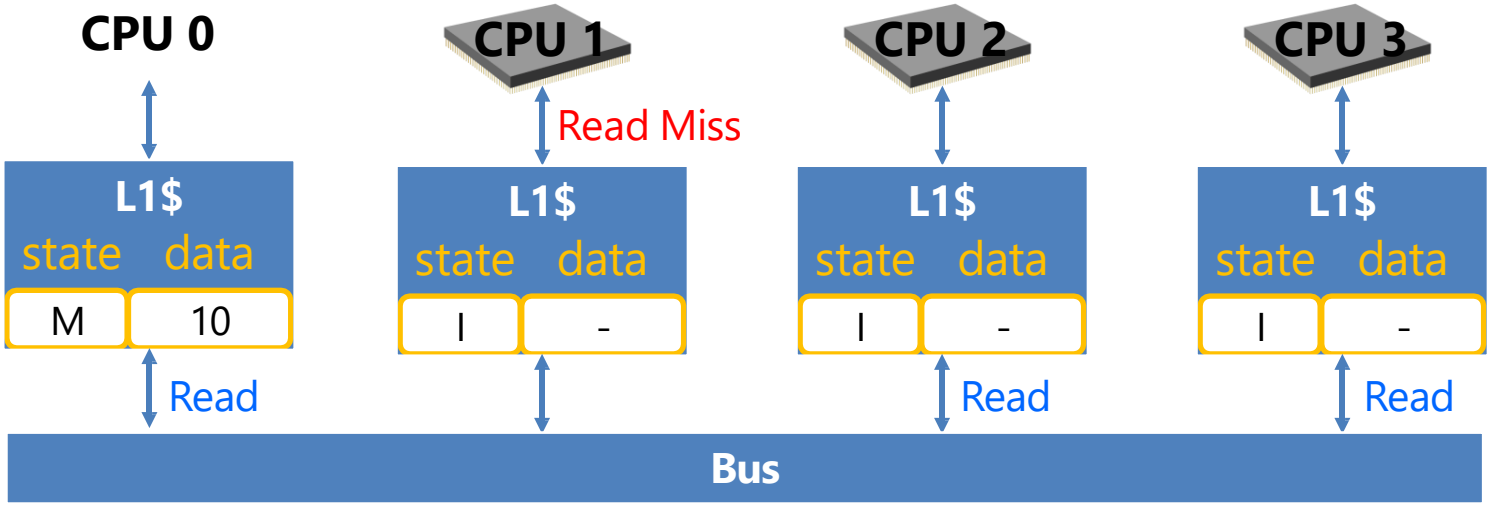

1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a Read Miss
3. A Read message is broadcast on the Bus

Scenario 1: Read Snooped on Bus

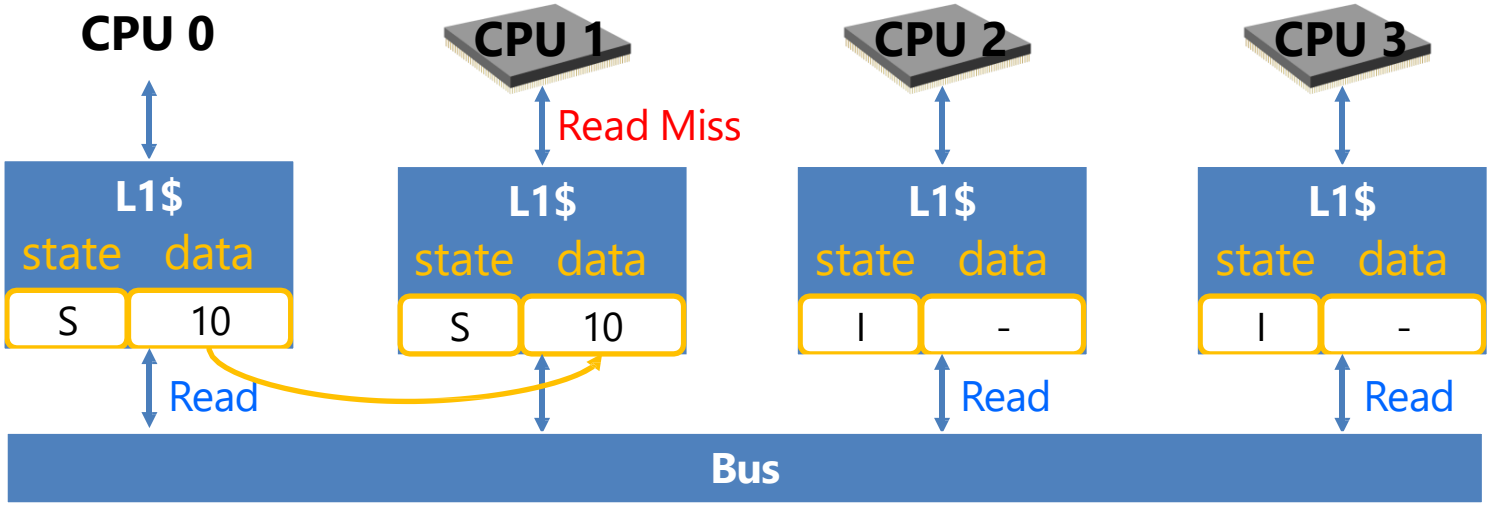

1. CPU 0: line in Modified state, CPU 1~3: line in Invalid state
2. CPU 1 reads line, causing a Read Miss
3. A Read message is broadcast on the Bus
4. CPU 0 snoops message and provides line to CPU1
   - Line in CPU 0 and CPU 1 both transition to Shared state

Scenario 1: Read Snooped on Bus

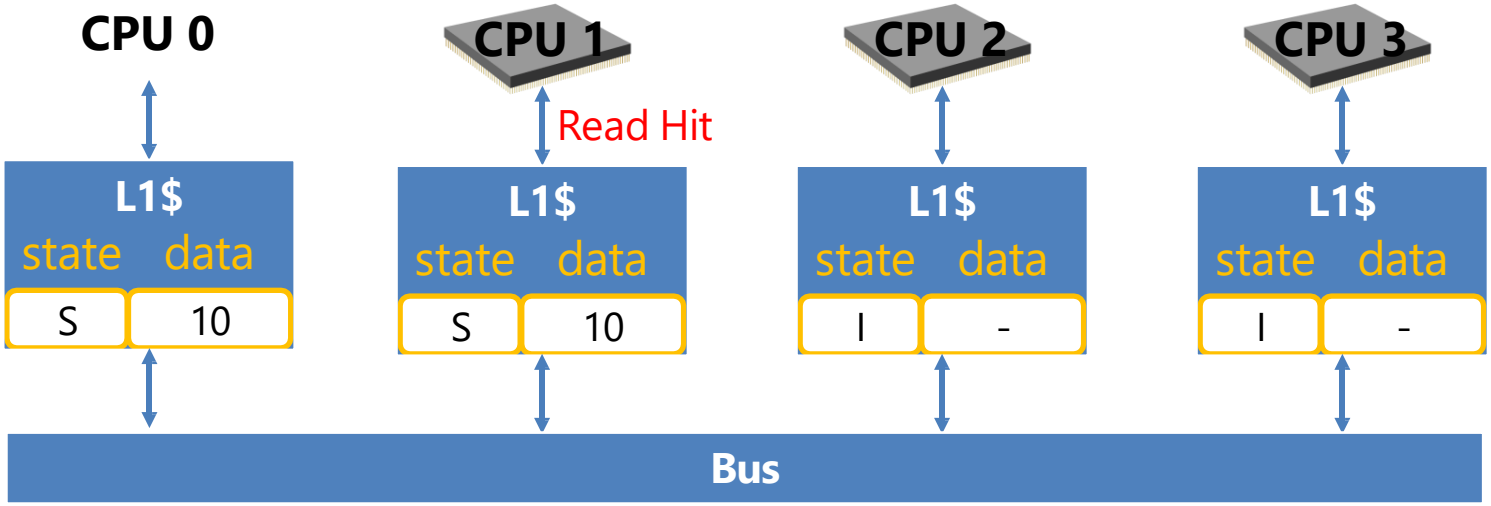

- Subsequent read hits on CPU 0 or CPU 1 don't generate messages
- Only read misses generate messages, not read hits
  → Reduces bus bandwidth pressure since misses are rare!

Scenario 2: Invalidate Snooped on Bus

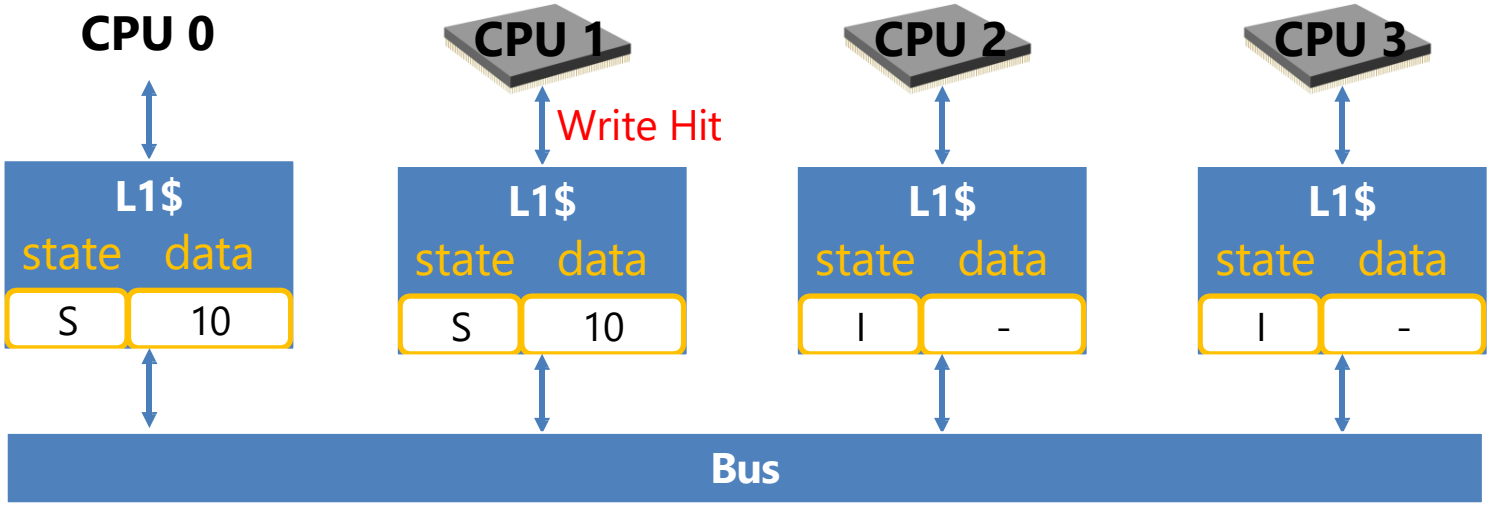

1. CPU 1 writes to line, causing a Write Hit

Scenario 2: Invalidate Snooped on Bus

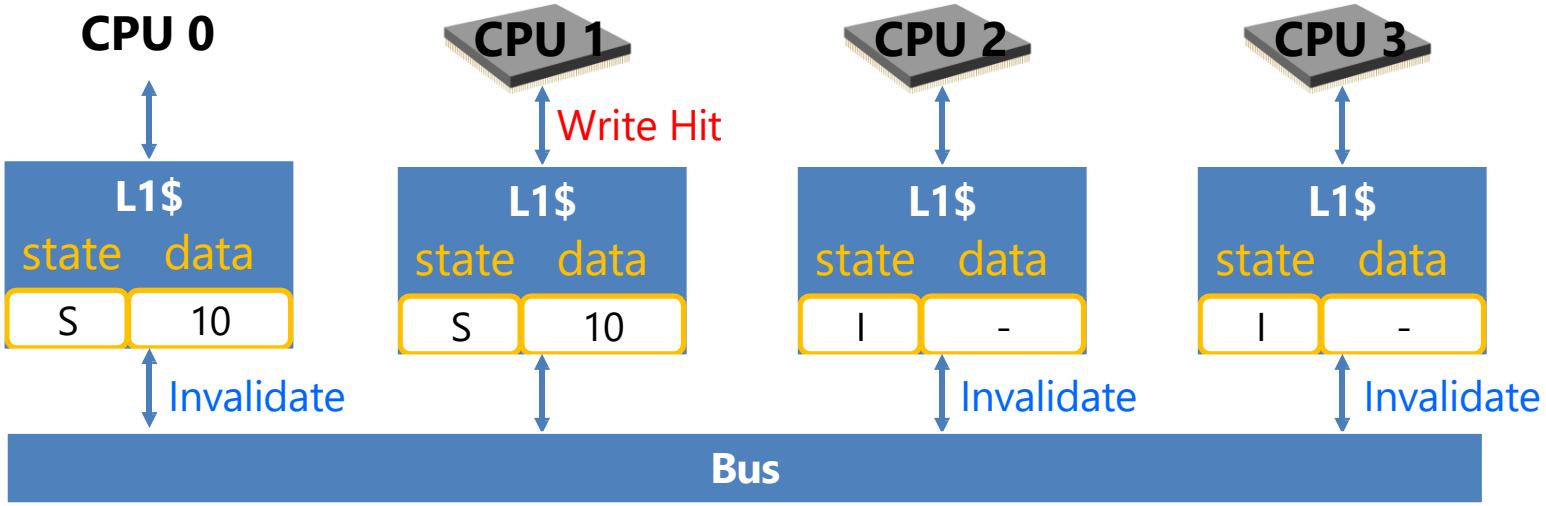

1. CPU 1 writes to line, causing a Write Hit
2. An Invalidate message is broadcast on the Bus
   - To remove all copies of the line that may become inconsistent

Scenario 2: Invalidate Snooped on Bus

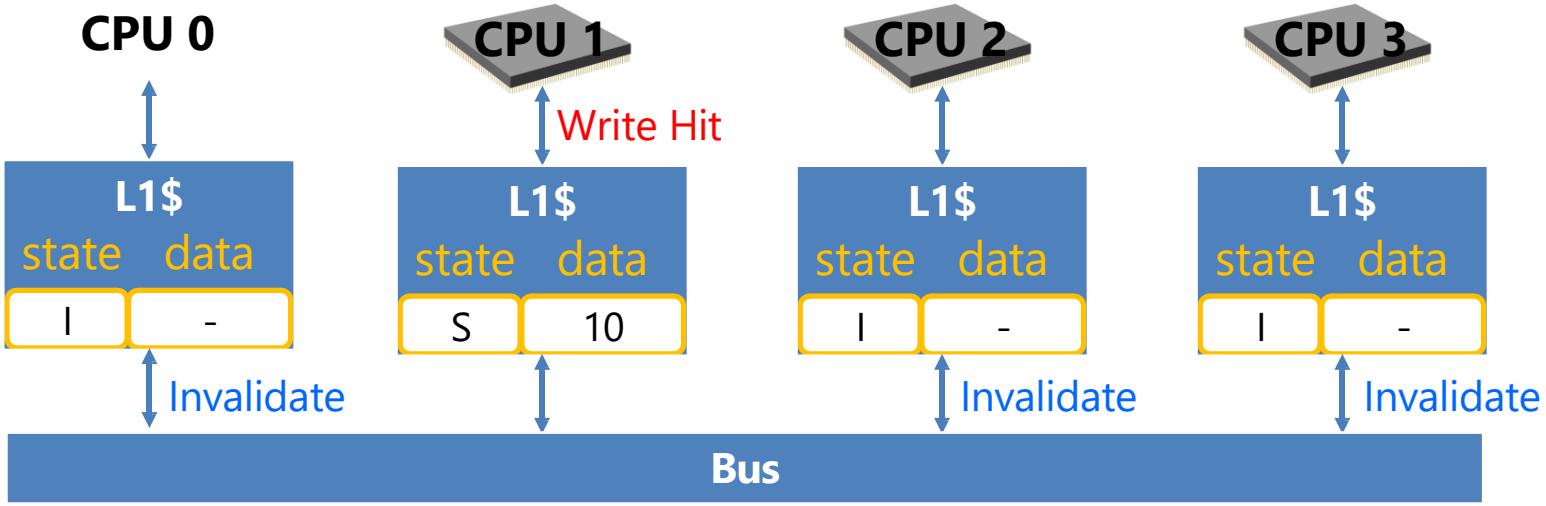

1. CPU 1 writes to line, causing a Write Hit
2. An Invalidate message is broadcast on the Bus
   - To remove all copies of the line that may become inconsistent
3. CPU 0 snoops message and invalidates its line

Scenario 2: Invalidate Snooped on Bus

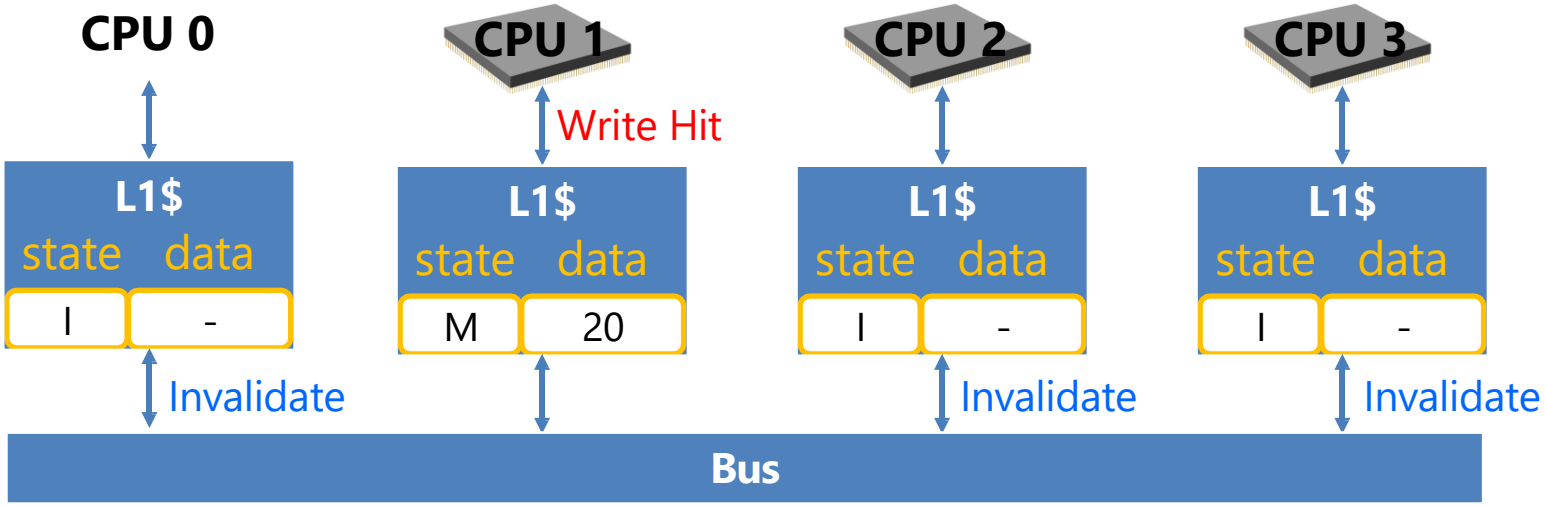

1. CPU 1 writes to line, causing a Write Hit
2. An Invalidate message is broadcast on the Bus
   - To remove all copies of the line that may become inconsistent
3. CPU 0 snoops message and invalidates its line
4. Line in CPU 1 transitions to Modified state

Scenario 2: Invalidate Snooped on Bus

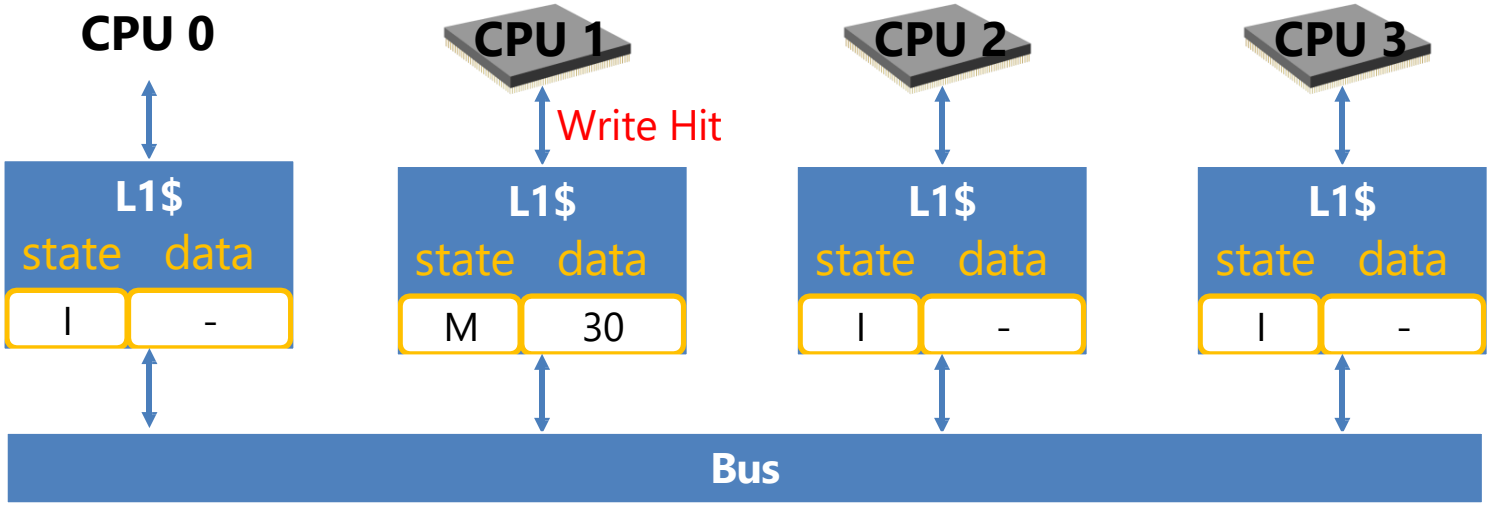

- Subsequent write hits on CPU 1 don't generate messages
- Only write hit in Shared state generate messages
  → Reduces bandwidth since most write hits are to Modified state

Scenario 3: Read Exclusive Snooped on Bus

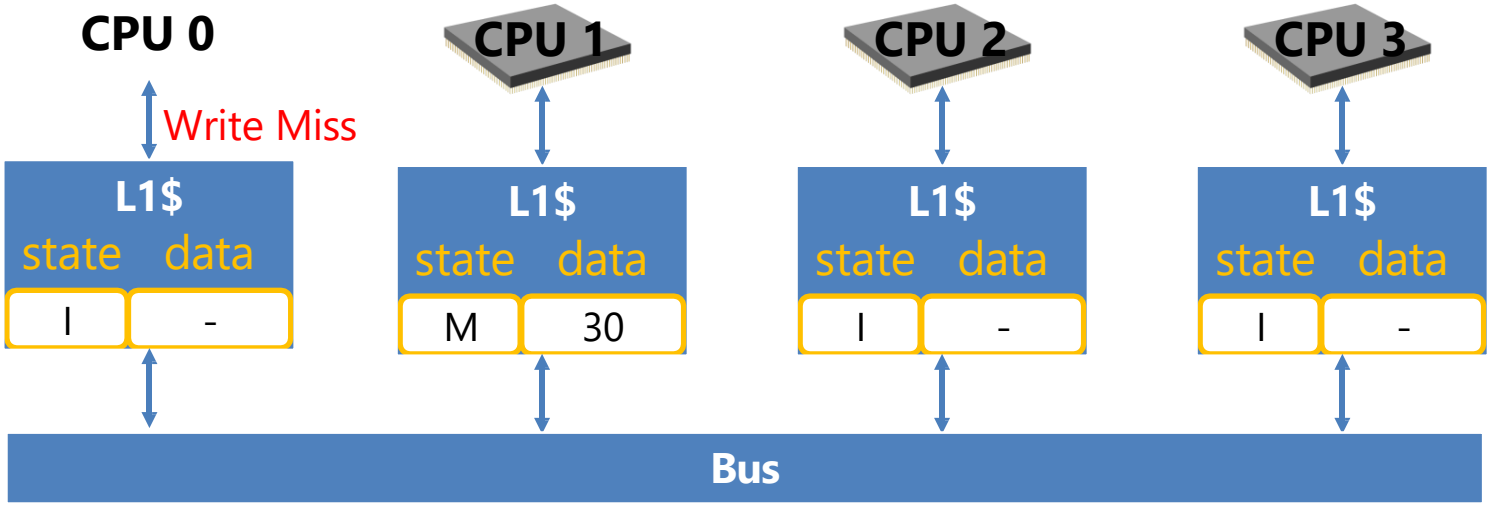

1. CPU 0 writes to line, causing a Write Miss

Scenario 3: Read Exclusive Snooped on Bus

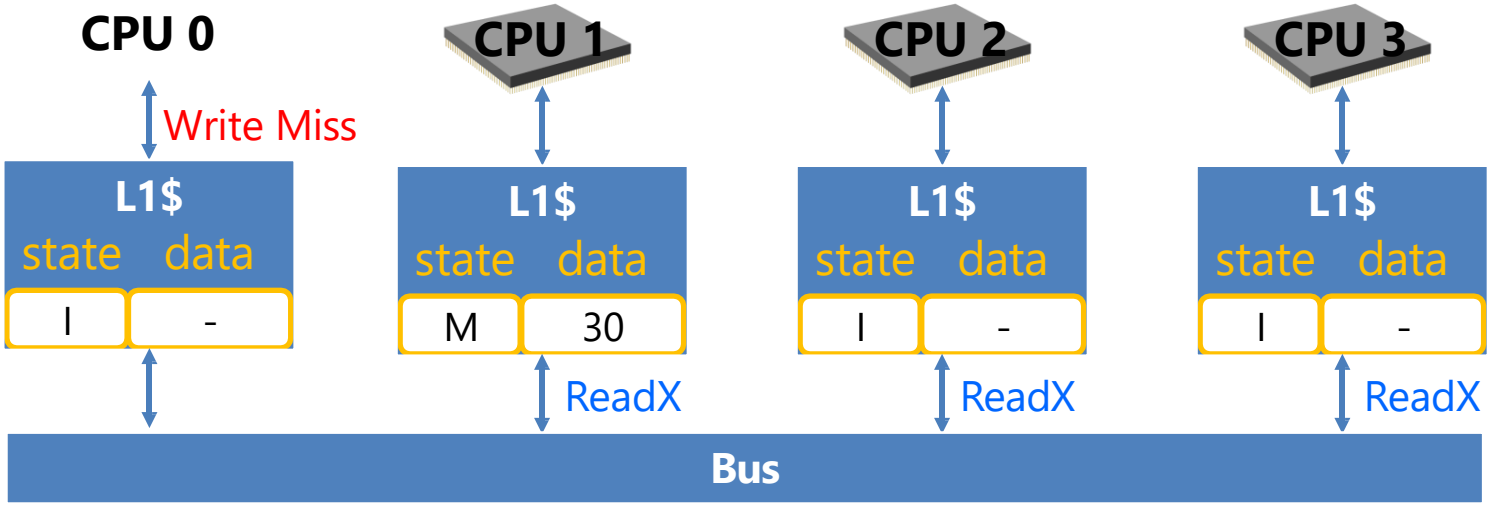

1. CPU 0 writes to line, causing a Write Miss
2. A ReadX (read exclusive) message is broadcast on the Bus
   - To read line into CPU 0 before writing to it (remember?)
   - To remove all copies of the line that may become inconsistent

Scenario 3: Read Exclusive Snooped on Bus

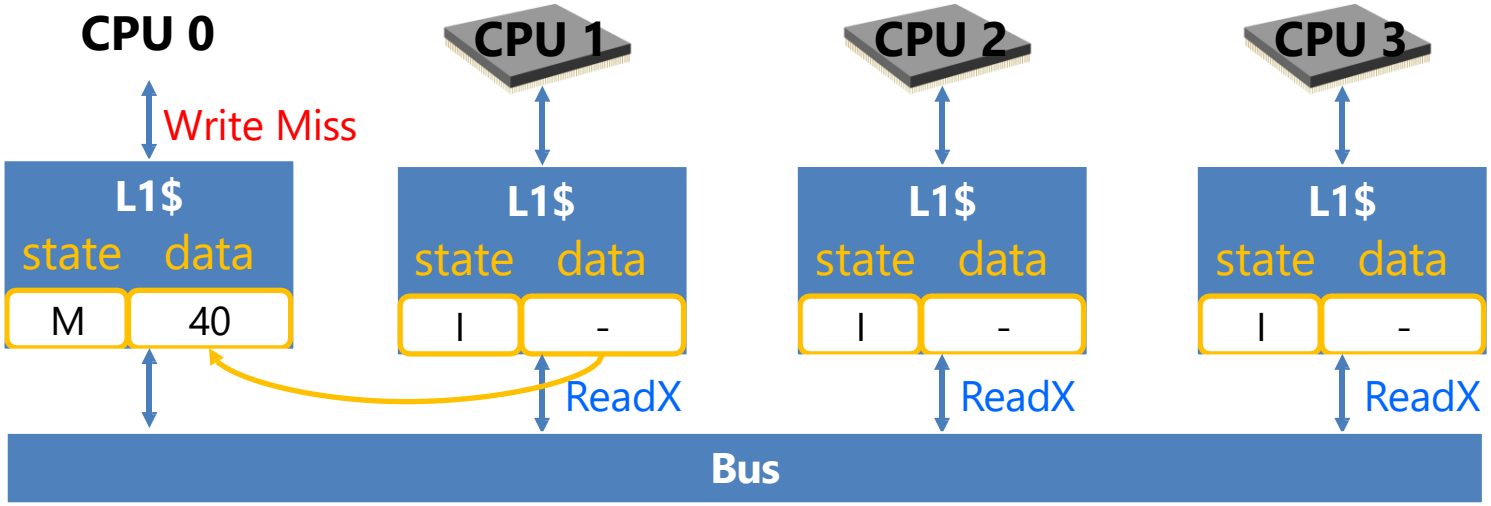

1. CPU 0 writes to line, causing a Write Miss
2. A ReadX (read exclusive) message is broadcast on the Bus
   - To read line into CPU 0 before writing to it (remember?)
   - To remove all copies of the line that may become inconsistent
3. CPU 1 snoops message and provides line to CPU 0
   - Line in CPU 1 is invalidated before line in CPU 0 is modified

Scenario 3: Read Exclusive Snooped on Bus

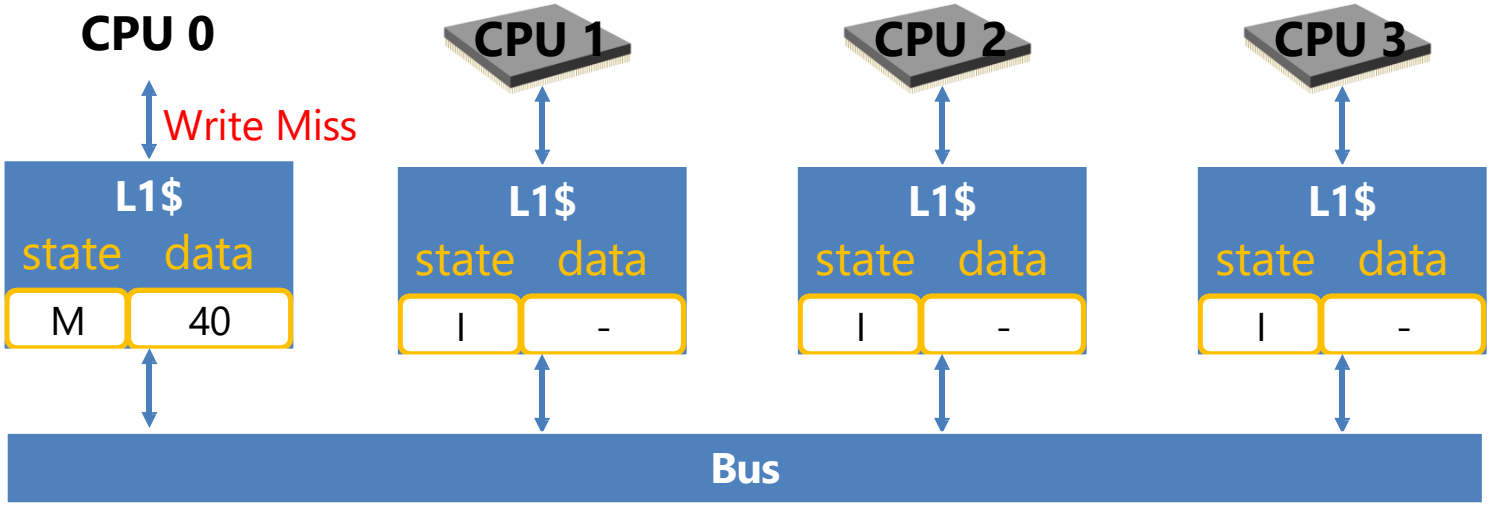

- Subsequent write hits on CPU 0 don't generate messages
- Write miss generates message, but not subsequent write hits
  → Reduces bus bandwidth pressure since misses are rare!

Cache Controller FSM for MSI Protocol

- Processor activity in red, Bus activity in blue

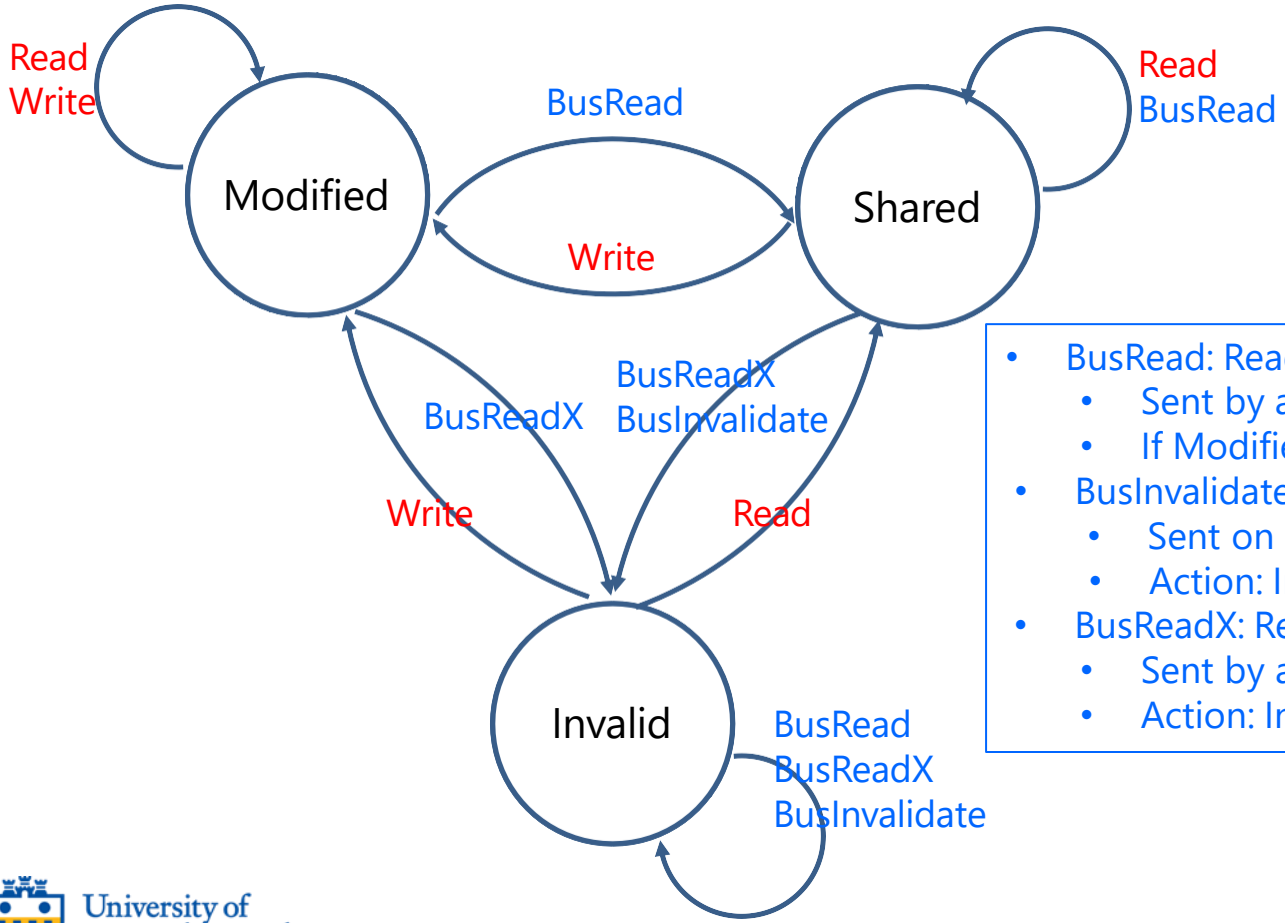

- BusRead: Read request is snooped
  - Sent by another processor on a read miss
  - If Modified, transition to Shared state
- BusInvalidate: Invalidate request is snooped
  - Sent on a write hit on shared cache line
  - Action: Invalidate shared line (if it exists)
- BusReadX: Read exclusive request is snooped
  - Sent by another processor on a write miss
  - Action: Invalidate shared line (if it exists)

TLB Coherence

University of Pittsburgh

How about TLBs?

- We said TLBs are also a type of cache that caches PTEs.
  - So what happens if a processor changes a PTE?
  - How does that change get propagated to other processor TLBs?

- Unfortunately, there is no hardware coherence for TLBs. ☹

- That means software (the OS) must handle the coherence
  - Which is of course much much slower

TLB shootdown

- In order to update a PTE (page table entry)
  - Initiator OS must first flush its own TLB
  - Send IPIs (Inter-processor interrupts) to other processors
    - To flush the TLBs for all other processors too
  - Source of significant performance overhead

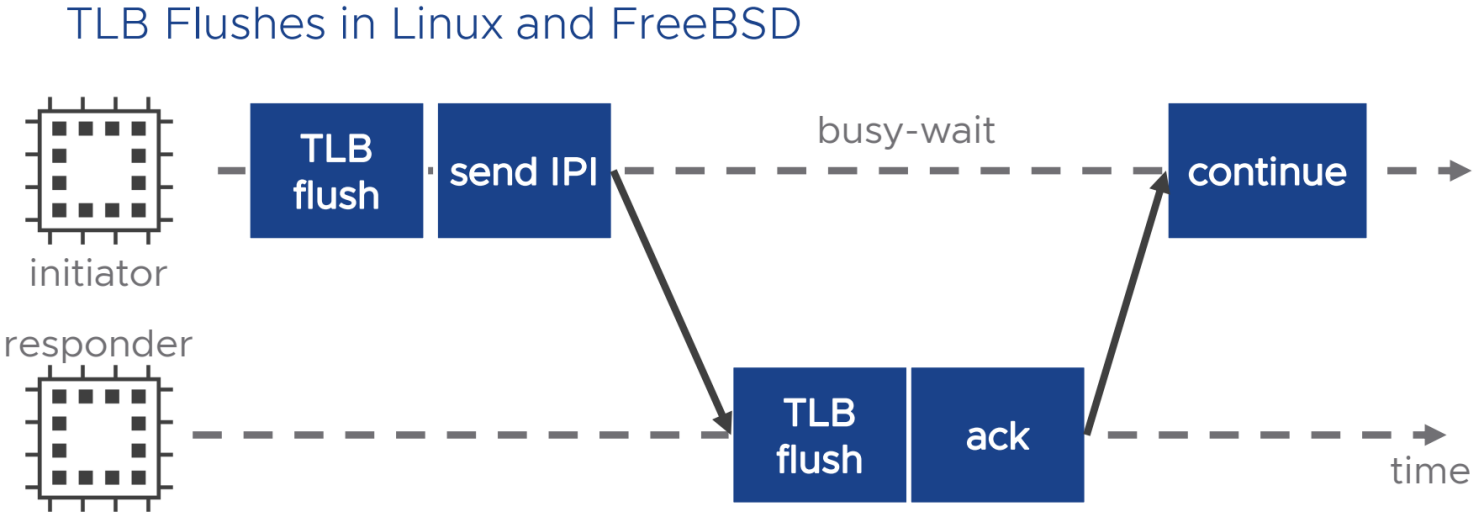

* Courtesy of Nadav Amit et al. at VMWare